

REFERENCE

Commands

Complete reference for commands available in Claude Code, including built-in commands and bundled skills.

Commands control Claude Code from inside a session. They provide a quick way to switch models, manage permissions, clear context, run a workflow, and more.

Type `/` to see every command available to you, or type `/` followed by letters to filter.

A command is only recognized at the start of your message. Text that follows the command name is passed to it as arguments.

Commands across a typical workflow

Most commands are useful at a specific point in a session, from setting up a project to shipping a change.

First session in a repo. Run `/init` to generate a starter `CLAUDE.md`, then `/memory` to refine it. Use `/mcp` and `/agents` to set up any servers or subagents the project needs, and `/permissions` to set the approval rules you want.

During a task. `/plan` switches into plan mode before a large change. `/model` and `/effort` adjust how much reasoning you're spending. When the conversation gets long, `/context` shows where the window is going and `/compact` summarizes it down; use `/btw` for a quick aside that shouldn't bloat history.

Running work in parallel. `/agents` opens the manager for the subagents Claude can delegate side tasks to, and `/tasks` lists what's running in the background of the current session. `/background` detaches the whole session to keep running as a background agent and frees your terminal. For a large change that spans the codebase, `/batch` decomposes it into independent units and runs each in its own worktree. See Run agents in parallel for how these approaches relate.

Before you ship. `/diff` shows what changed, `/code-review` checks the diff for correctness bugs and cleanups and can apply the findings with `--fix`, `/review` runs the same read-only review on a

GitHub pull request, and `/security-review` gives a deeper read-only pass. `/code-review ultra` runs a multi-agent review in the cloud.

Between sessions. `/clear` starts fresh on a new task while keeping project memory. `/resume` and `/branch` let you return to or fork an earlier conversation. `/teleport` pulls a web session into this terminal, and `/remote-control` lets you continue this local session from another device.

When something is wrong. `/rewind` rolls code and conversation back to a checkpoint, or summarizes part of the conversation. `/doctor` and `/debug` diagnose install and runtime issues, and `/feedback` reports a bug with session context attached.

All commands

The table below lists all the commands included in Claude Code. Most are built-in commands whose behavior is coded into the CLI. Two kinds of entries are marked:

- **Skill**: a bundled skill. It works like skills you write yourself: a prompt handed to Claude, which Claude can also invoke automatically when relevant.
- **Workflow**: a bundled dynamic workflow that fans work out across many subagents and runs in the background.

To add your own commands, see skills.

In the table below, `<arg>` indicates a required argument and `[arg]` indicates an optional one.

ⓘ Not every command appears for every user. Availability depends on your platform, plan, and environment. For example, `/desktop` only shows on macOS and Windows when signed in with a Claude subscription, and `/upgrade` only shows on Pro and Max plans.

Command	Purpose
<code>/add-dir <path></code>	Add a working directory for file access during the current session. Most <code>.claude/</code> configuration is <u>not discovered</u> from the added directory. You can later resume the session from the added directory with <code>--continue</code> or <code>--resume</code>
<code>/advisor [model off]</code>	Enable or disable the <u>advisor tool</u> , which consults a second model for guidance at key moments during a task. Accepts <code>opus</code> , <code>sonnet</code> , <code>fable</code> (v2.1.170+), or a full model ID. Without an argument, opens a picker. Requires Claude Code v2.1.98 or later
<code>/agents</code>	Manage <u>agent</u> configurations
<code>/autofix-pr [prompt]</code>	Spawn a <u>Claude Code on the web</u> session that watches the current branch's PR and pushes fixes when CI fails or reviewers leave comments. Detects the open PR from your checked-out branch with <code>gh pr view</code> ; to watch a different PR, check out its

Command	Purpose
	branch first. By default the cloud session is told to fix every CI failure and review comment; pass a prompt to give it different instructions, for example <code>/autofix-pr only fix lint and type errors</code> . Requires the <code>gh</code> CLI and access to <u>Claude Code on the web</u>
<code>/background [prompt]</code>	Detach the current session to run as a <u>background agent</u> and free this terminal. Pass a prompt to send one more instruction before detaching. Monitor the session with <code>claude agents</code> . Alias: <code>/bg</code>
<code>/batch <instruction></code>	Skill. Orchestrate large-scale changes across a codebase in parallel. Researches the codebase, decomposes the work into 5 to 30 independent units, and presents a plan. Once approved, spawns one <u>background subagent</u> per unit in an isolated <u>git worktree</u> . Each subagent implements its unit, runs tests, and opens a pull request. Requires a git repository. Example: <code>/batch migrate src/ from Solid to React</code>
<code>/branch [name]</code>	Create a branch of the current conversation at this point, so you can try a different direction without losing the conversation as it stands. Switches you into the branch and preserves the original, which you can return to with <code>/resume</code> . To hand a side task to a background subagent instead of switching into a copy yourself, use <code>/fork</code>
<code>/btw <question></code>	Ask a quick <u>side question</u> without adding to the conversation
<code>/cd <path></code>	Move this session to a new working directory. The conversation's prompt cache is preserved: the new directory's <u>CLAUDE.md</u> is appended as a message instead of rebuilding the system prompt. The session is relocated to the new directory's project storage, so <code>--resume</code> and <code>--continue</code> find it from there. Prompts you to trust the directory if you haven't worked in it before. To grant access to an extra directory without moving the session, use <code>/add-dir</code> . Restrict or disable <code>/cd</code> targets with <u>cd permission rules</u> . Requires Claude Code v2.1.169 or later; earlier versions report <code>Unknown command: /cd</code>
<code>/chrome</code>	Configure <u>Claude in Chrome</u> settings
<code>/claude-api</code> <code>[migrate managed-agents-onboard]</code>	Skill. Load Claude API reference material for your project's language (Python, TypeScript, Java, Go, Ruby, C#, PHP, or cURL) and Managed Agents reference. Covers tool use, streaming, batches, structured outputs, and common pitfalls. Also activates automatically when your code imports <code>anthropic</code> or <code>@anthropic-ai/sdk</code> . Run <code>/claude-api migrate</code> to upgrade existing Claude API code to a newer model: Claude asks which files to scan and which model to target, then updates model IDs, thinking configuration, and other parameters that changed between versions. Run <code>/claude-api managed-agents-onboard</code> for an interactive walkthrough that creates a new Managed Agent from scratch
<code>/clear [name]</code>	Start a new conversation with empty context. The previous conversation stays available in <code>/resume</code> . Pass a name to label the previous conversation in the <code>/resume</code> picker. To free up context while continuing the same conversation, use <code>/compact</code> instead. Aliases: <code>/reset</code> , <code>/new</code>
<code>/code-review</code> <code>[low medium high xhigh max ultra] [--fix] [--comment] [target]</code>	Skill. Review the current diff for correctness bugs and for reuse, simplification, and efficiency cleanups. Pass <code>--fix</code> to apply findings to your working tree, <code>--comment</code> to post them as inline GitHub PR comments, or <code>ultra</code> to run a deep <u>cloud review</u> . From v2.1.154, <code>/simplify</code> runs a separate cleanup-only review that applies fixes without hunting for bugs. See <u>Review a diff locally</u> for effort levels and targeting

Command	Purpose
<code>/color [color default]</code>	Set the prompt bar color for the current session. Available colors: <code>red</code> , <code>blue</code> , <code>green</code> , <code>yellow</code> , <code>purple</code> , <code>orange</code> , <code>pink</code> , <code>cyan</code> . Use <code>default</code> to reset, or run with no argument to pick a random color. When <u>Remote Control</u> is connected, the color syncs to <code>claude.ai/code</code>
<code>/compact</code> <code>[instructions]</code>	Free up context by summarizing the conversation so far. Optionally pass focus instructions for the summary. See <u>how compaction handles rules, skills, and memory files</u>
<code>/config [key=value ...]</code>	Open the <u>Settings</u> interface to adjust theme, model, <u>output style</u> , and other preferences. From v2.1.181, pass one or more <code>key=value</code> pairs to set a setting directly without opening the interface, for example <code>/config thinking=false</code> . From v2.1.182, named shorthand keys are also accepted, such as <code>/config theme=dark</code> or <code>/config model=sonnet</code> . The <code>key=value</code> form also works in non-interactive mode (<code>-p</code>) and from <u>Remote Control</u> . Run <code>/config --help</code> to list every settable key with its options. Alias: <code>/settings</code>
<code>/context [all]</code>	Visualize current context usage as a colored grid. Shows optimization suggestions for context-heavy tools, memory bloat, and capacity warnings. In <u>fullscreen mode</u> the per-item breakdown is collapsed to keep the grid visible. Pass <code>all</code> to expand it
<code>/copy [N]</code>	Copy the last assistant response to clipboard. Pass a number <code>N</code> to copy the Nth-latest response: <code>/copy 2</code> copies the second-to-last. When code blocks are present, shows an interactive picker to select individual blocks or the full response. Press <code>w</code> in the picker to write the selection to a file instead of the clipboard, which is useful over SSH
<code>/cost</code>	Alias for <code>/usage</code>
<code>/debug [description]</code>	<u>Skill.</u> Enable debug logging for the current session and troubleshoot issues by reading the session debug log. Debug logging is off by default unless you started with <code>claude --debug</code> , so running <code>/debug</code> mid-session starts capturing logs from that point forward. Optionally describe the issue to focus the analysis
<code>/deep-research</code> <code><question></code>	<u>Workflow.</u> Fan out web searches on a question, fetch and cross-check sources, and synthesize a cited report
<code>/desktop</code>	Continue the current session in the Claude Code Desktop app. Requires macOS or Windows and a Claude subscription. Alias: <code>/app</code>
<code>/diff</code>	Open an interactive diff viewer showing uncommitted changes and per-turn diffs. Use left/right arrows to switch between the current git diff and individual Claude turns, and up/down to browse files
<code>/doctor</code>	Diagnose and verify your Claude Code installation and settings. Results show with status icons. Press <code>f</code> to have Claude fix any reported issues

Command	Purpose
<code>/effort [level auto]</code>	Set the model effort level . Accepts <code>low</code> , <code>medium</code> , <code>high</code> , <code>xhigh</code> , <code>max</code> , or <code>ultracode</code> ; available levels depend on the model, and <code>max</code> and <code>ultracode</code> are session-only. <code>ultracode</code> is a Claude Code setting that combines <code>xhigh</code> reasoning with automatic workflow orchestration. <code>auto</code> resets to the model default. Without an argument, opens an interactive slider; use left and right arrows to pick a level and <code>Enter</code> to apply. Takes effect immediately without waiting for the current response to finish
<code>/exit</code>	Exit the CLI. In an attached background session , this detaches and the session keeps running. Alias: <code>/quit</code>
<code>/export [filename]</code>	Export the current conversation as plain text. With a filename, writes directly to that file. Without, opens a dialog to copy to clipboard or save to a file
<code>/fast [on off]</code>	Toggle fast mode on or off
<code>/feedback [report]</code>	Submit feedback, report a bug, or share your conversation. Aliases: <code>/bug</code> , <code>/share</code>
<code>/fewer-permission-prompts</code>	Skill . Scan your transcripts for common read-only Bash and MCP tool calls, then add a prioritized allowlist to project <code>.claude/settings.json</code> to reduce permission prompts
<code>/focus</code>	Toggle the focus view, which shows only your last prompt, a one-line tool-call summary with edit diffstats, and the final response. The selection persists across sessions; set <code>viewMode</code> in settings to override it. Only available in fullscreen rendering
<code>/fork <directive></code>	Spawn a forked subagent : a background subagent that inherits the full conversation and works on the directive while you keep going. Its result returns to your conversation when it finishes. To switch into a copy of the conversation yourself, use <code>/branch</code> . Before v2.1.161, <code>/fork</code> is an alias for <code>/branch</code>
<code>/goal [condition clear]</code>	Set a goal : Claude keeps working across turns until the condition is met. With no argument, shows the current or most recently achieved goal. <code>clear</code> , <code>stop</code> , <code>off</code> , <code>reset</code> , <code>none</code> , or <code>cancel</code> removes an active goal early
<code>/heapdump</code>	Write a JavaScript heap snapshot and a memory breakdown to <code>~/Desktop</code> , or your home directory on Linux without a Desktop folder, for diagnosing high memory usage. See troubleshooting
<code>/help</code>	Show help and available commands
<code>/hooks</code>	View hook configurations for tool events
<code>/ide</code>	Manage IDE integrations and show status
<code>/init</code>	Initialize project with a <code>CLAUDE.md</code> guide. Set <code>CLAUDE_CODE_NEW_INIT=1</code> for an interactive flow that also walks through skills, hooks, and personal memory files
<code>/insights</code>	Generate a report analyzing your Claude Code sessions, including project areas, interaction patterns, and friction points
<code>/install-github-app</code>	Install the Claude GitHub App for a repository, with an optional step to set up GitHub Actions workflows and secrets. Walks you through selecting a repo and configuring

Command	Purpose
	the integration
<code>/install-slack-app</code>	Install the Claude Slack app. Opens a browser to complete the OAuth flow
<code>/keybindings</code>	Open your <u>keyboard shortcuts</u> file
<code>/login</code>	Sign in to your Anthropic account
<code>/logout</code>	Sign out from your Anthropic account
<code>/loop [interval]</code> <code>[prompt]</code>	Skill. Run a prompt repeatedly while the session stays open. Omit the interval and Claude self-paces between iterations. Omit the prompt and, where available , Claude runs an autonomous maintenance check or the prompt in <code>.claude/loop.md</code> . Example: <code>/loop 5m check if the deploy finished</code> . See <u>Run prompts on a schedule</u> . Alias: <code>/proactive</code>
<code>/mcp [reconnect</code> <code><server> enable disable</code> <code>[<server> all]]</code>	Manage MCP server connections and OAuth authentication. Run with no argument to open the interactive list, pass <code>reconnect <server></code> to reconnect one disconnected server, or pass <code>enable / disable</code> with a server name or <code>all</code> to change connection state without opening the dialog
<code>/memory</code>	Edit <code>CLAUDE.md</code> memory files, enable or disable <u>auto-memory</u> , and view auto-memory entries
<code>/mobile</code>	Show QR code to download the Claude mobile app. Aliases: <code>/ios</code> , <code>/android</code>
<code>/model [model]</code>	Switch the AI model and save it as your default for new sessions. For models that support it, use left/right arrows to <u>adjust effort level</u> . With no argument, opens a picker; press <code>s</code> on a row to switch for the current session only. The picker asks for confirmation when the conversation has prior output, since the next response re-reads the full history without cached context. Once confirmed, the change applies without waiting for the current response to finish
<code>/passes</code>	Share a free week of Claude Code with friends. Only visible if your account is eligible
<code>/permissions</code>	Manage allow, ask, and deny rules for tool permissions. Opens an interactive dialog where you can view rules by scope, add or remove rules, manage working directories, and review <u>recent auto mode denials</u> . Alias: <code>/allowed-tools</code>
<code>/plan [description]</code>	Enter plan mode directly from the prompt. Pass an optional description to enter plan mode and immediately start with that task, for example <code>/plan fix the auth bug</code>
<code>/plugin [subcommand]</code>	Manage Claude Code <u>plugins</u> . Run with no argument to open the plugin menu, or pass a subcommand such as <code>list</code> , <code>install</code> , <code>enable</code> ,or <code>disable</code> to act directly
<code>/powerup</code>	Discover Claude Code features through quick interactive lessons with animated demos
<code>/pr-comments [PR]</code>	Removed in v2.1.91. Ask Claude directly to view pull request comments instead. On earlier versions, fetches and displays comments from a GitHub pull request; automatically detects the PR for the current branch, or pass a PR URL or number. Requires the <code>gh</code> CLI

Command	Purpose
<code>/privacy-settings</code>	View and update your privacy settings. Only available for Pro and Max plan subscribers
<code>/radio</code>	Open Claude FM lo-fi radio in your browser. Prints the stream URL when no browser is available. Not available on Bedrock, Vertex, or Foundry
<code>/recap</code>	Generate a one-line summary of the current session on demand. See Session recap for the automatic recap that appears after you've been away
<code>/release-notes</code>	View the changelog in an interactive version picker. Select a specific version to see its release notes, or choose to show all versions
<code>/reload-plugins [--force]</code>	Reload all active plugins to apply pending changes without restarting. Reports counts for each reloaded component and flags any load errors. When the reload would change which MCP tools are loaded and invalidate the prompt cache, the command warns and skips unless you pass <code>--force</code>
<code>/reload-skills</code>	Re-scan skill and command directories so skills added or changed on disk during the session become available without restarting. Reports how many skills are available and how many were added or removed. Added in v2.1.152
<code>/remote-control</code>	Make this session available for remote control from claude.ai. Alias: <code>/rc</code>
<code>/remote-env</code>	Choose the default environment for cloud agents
<code>/rename [name]</code>	Rename the current session and show the name on the prompt bar. Without a name, auto-generates one from conversation history
<code>/resume [session]</code>	Resume a conversation by ID or name, or open the session picker. As of v2.1.144, background sessions appear in the picker marked with <code>bg</code> . Alias: <code>/continue</code>
<code>/review [PR]</code>	Review a GitHub pull request by number, using the same review engine as <code>/code-review</code> . With no arguments, lists open PRs to pick from. For a cloud-based review, see /code-review ultra
<code>/rewind</code>	Rewind the conversation and/or code to a previous point, or summarize from a selected message. See checkpointing . Aliases: <code>/checkpoint</code> , <code>/undo</code>
<code>/run</code>	Skill. Launch and drive your project's app to see a change working in the running app, not just in tests. See Run and verify your app . Requires Claude Code v2.1.145 or later
<code>/run-skill-generator</code>	Skill. Teach <code>/run</code> and <code>/verify</code> how to build, launch, and drive your project's app from a clean environment by writing a per-project skill . Requires Claude Code v2.1.145 or later
<code>/sandbox</code>	Toggle sandbox mode . Available on supported platforms only
<code>/schedule [description]</code>	Create, update, list, or run routines , which execute on Anthropic-managed cloud infrastructure. Claude walks you through the setup conversationally. Alias: <code>/routines</code>
<code>/scroll-speed</code>	Adjust mouse wheel scroll speed interactively, with a ruler you can scroll while the dialog is open to preview the change. Available in fullscreen rendering only and not in the JetBrains IDE terminal

Command	Purpose
<code>/security-review</code>	Analyze pending changes on the current branch for security vulnerabilities. Reviews the git diff and identifies risks like injection, auth issues, and data exposure
<code>/setup-bedrock</code>	Configure Amazon Bedrock authentication, region, and model pins through an interactive wizard. Only visible when <code>CLAUDE_CODE_USE_BEDROCK=1</code> is set. First-time Bedrock users can also access this wizard from the login screen
<code>/setup-vertex</code>	Configure Google Vertex AI authentication, project, region, and model pins through an interactive wizard. Only visible when <code>CLAUDE_CODE_USE_VERTEX=1</code> is set. First-time Vertex AI users can also access this wizard from the login screen
<code>/simplify [target]</code>	Skill. Review the changed code for cleanup opportunities and apply the fixes. Four review agents run in parallel, covering reuse of existing helpers, simplification, efficiency, and whether the change sits at the right level of abstraction. From v2.1.154, the review does not look for correctness bugs. Use <code>/code-review</code> to find bugs. On earlier versions <code>/simplify</code> is equivalent to <code>/code-review --fix</code> . Pass a path or PR reference to review a specific target
<code>/skills</code>	List available skills . Press <code>t</code> to sort by token count. Press <code>Space</code> to hide a skill from Claude or the / menu , then <code>Enter</code> to save
<code>/stats</code>	Alias for <code>/usage</code> . Opens on the Stats tab
<code>/status</code>	Open the Settings interface (Status tab) showing version, model, account, and connectivity. Works while Claude is responding, without waiting for the current response to finish
<code>/statusline</code>	Configure Claude Code's status line . Describe what you want, or run without arguments to auto-configure from your shell prompt
<code>/stickers</code>	Order Claude Code stickers
<code>/stop</code>	Stop the current background session . Only available while attached to a background session; the transcript and any worktree are kept. To detach without stopping, use <code>/exit</code> or press <code>←</code>
<code>/tasks</code>	View and manage everything running in the background. Also available as <code>/bashes</code>
<code>/team-onboarding</code>	Generate a team onboarding guide from your Claude Code usage history. Claude analyzes your sessions, commands, and MCP server usage from the past 30 days and produces a markdown guide a teammate can paste as a first message to get set up quickly. For claude.ai subscribers on Pro, Max, Team, and Enterprise plans, also returns a share link teammates can open directly in Claude Code
<code>/teleport</code>	Pull a Claude Code on the web session into this terminal: opens a picker, then fetches the branch and conversation. Also available as <code>/tp</code> . Requires a claude.ai subscription
<code>/terminal-setup</code>	Configure terminal keybindings for Shift+Enter and other shortcuts. Only visible in terminals that need it, like VS Code, Cursor, Devin Desktop, Alacritty, or Zed
<code>/theme</code>	Change the color theme. Includes an <code>auto</code> option that matches your terminal's light or dark background, light and dark variants, colorblind-accessible (daltonized) themes, ANSI themes that use your terminal's color palette, and any custom themes

Command	Purpose
	from <code> ~/.claude/themes/ </code> or plugins. Select New custom theme... to create one
<code>/tui</code> <code>[default fullscreen]</code>	Set the terminal UI renderer and relaunch into it with your conversation intact. <code>fullscreen</code> enables the <u>flicker-free alt-screen renderer</u> . With no argument, prints the active renderer
<code>/ultraplan <prompt></code>	Draft a plan in an <u>ultraplan</u> session, review it in your browser, then execute remotely or send it back to your terminal
<code>/ultrareview [PR]</code>	Run a deep, multi-agent code review in a cloud sandbox with <u>ultrareview</u> . The preferred invocation is now <code>/code-review ultra</code> , and <code>/ultrareview</code> remains as an alias. Includes 3 free runs on Pro and Max, then requires <u>usage credits</u>
<code>/upgrade</code>	Open the upgrade page to switch to a higher plan tier
<code>/usage</code>	Show session cost, plan usage limits, and activity stats. On a Pro, Max, Team, or Enterprise plan, includes a breakdown of usage by skill, subagent, plugin, and MCP server. See the <u>cost tracking guide</u> for details. <code>/cost</code> and <code>/stats</code> are aliases
<code>/usage-credits</code>	Configure usage credits to keep working when you hit a limit. Previously <code>/extra-usage</code>
<code>/verify</code>	Skill. Confirm a code change does what it should by building your project's app, running it, and observing the result, rather than relying on tests or type checks. See <u>Run and verify your app</u> . Requires Claude Code v2.1.145 or later
<code>/vim</code>	Removed in v2.1.92. To toggle between Vim and Normal editing modes, use <code>/config</code> → Editor mode
<code>/voice [hold tap off]</code>	Toggle <u>voice dictation</u> , or enable it in a specific mode. Requires a Claude.ai account
<code>/web-setup</code>	Connect your GitHub account to <u>Claude Code on the web</u> using your local <code>gh</code> CLI credentials. <code>/schedule</code> prompts for this automatically if GitHub isn't connected
<code>/workflows</code>	Open the <u>workflow</u> progress view to watch, pause, resume, or save running and completed workflows

MCP prompts

MCP servers can expose prompts that appear as commands. These use the format

`/mcp__<server>__<prompt>` and are dynamically discovered from connected servers. See MCP prompts for details.

See also

- Skills: create your own commands
- Interactive mode: keyboard shortcuts, Vim mode, and command history

- CLI reference: launch-time flags