

REFERENCE

Environment variables

Reference for environment variables that control Claude Code behavior.

Environment variables can control Claude Code behavior such as model selection, authentication, request routing, and feature toggles. Many of the same behaviors can also be configured through a **settings file** field, a **CLI flag**, or an in-session command like `/model` .

This page covers how to:

- **Set environment variables** in your shell or in a settings file
- **Check which value applies** when a behavior can be set more than one way
- **Look up the variables Claude Code reads**

Set environment variables

A variable you set in your shell lasts for that terminal session, while a variable in a settings file applies every time `claude` runs.

In your shell

Set the variable before launching `claude` :

macOS, Linux, WSL

```
export API_TIMEOUT_MS="1200000"
claude
```

To set it for every session, add the `export` line to `~/.bashrc` , `~/.zshrc` , or your shell's profile file.

Windows PowerShell

```
$env:API_TIMEOUT_MS = "1200000"
claude
```

To set it for every session, run `[Environment]::SetEnvironmentVariable("API_TIMEOUT_MS", "1200000", "User")` and open a new terminal.

Windows CMD

```
set API_TIMEOUT_MS=1200000
claude
```

To set it for every session, run `setx API_TIMEOUT_MS "1200000"` and open a new terminal.

In settings files

Add variables under the `env` key in a `settings.json` file. Claude Code reads them directly from the file at startup, so they take effect no matter how `claude` was launched.

~/claude/settings.json

```
{
  "env": {
    "API_TIMEOUT_MS": "1200000",
    "BASH_DEFAULT_TIMEOUT_MS": "300000"
  }
}
```

The file you choose controls who the variables apply to:

File	Applies to
~/claude/settings.json	You, in every project
.claude/settings.json	Everyone working in the project, checked into source control

File	Applies to
<code>.claude/settings.local.json</code>	You, in this project only (add it to your gitignore if you create it by hand)
Managed settings	Everyone in your organization, deployed by an admin

See [Settings files](#) for where each file lives and [Settings precedence](#) for how they combine when more than one sets the same variable.

Precedence

Where the same behavior has both an environment variable and a settings field, the environment variable takes precedence. For example, `ANTHROPIC_MODEL` overrides the `model` setting, and `CLAUDE_CODE_AUTO_CONNECT_IDE` overrides `autoConnectIde`. The settings field applies when the environment variable is not set.

How an environment variable interacts with CLI flags and in-session commands varies per feature: `-model` and `/model` override `ANTHROPIC_MODEL`, while `CLAUDE_CODE_EFFORT_LEVEL` overrides `/effort`. When a variable interacts with another configuration source, its row in the [Variables](#) list states the precedence or links to the page that documents it.

Claude Code reads environment variables at startup, so changes take effect the next time you launch `claude`.

Variables

Variable	Purpose
<code>ANTHROPIC_API_KEY</code>	API key sent as <code>X-API-Key</code> header. When set, this key is used instead of your Claude Pro, Max, Team, or Enterprise subscription even if you are logged in. In non-interactive mode (<code>-p</code>), the key is always used when present. In interactive mode, you are prompted to approve the key once before it overrides your subscription. To use your subscription instead, run <code>unset ANTHROPIC_API_KEY</code>
<code>ANTHROPIC_AUTH_TOKEN</code>	Custom value for the <code>Authorization</code> header (the value you set here will be prefixed with <code>Bearer</code>)
<code>ANTHROPIC_AWS_API_KEY</code>	Workspace API key for Claude Platform on AWS , generated in the AWS Console. Sent as <code>x-api-key</code> and takes precedence over AWS SigV4
<code>ANTHROPIC_AWS_BASE_URL</code>	Override the Claude Platform on AWS endpoint URL. Use for custom regions or when routing through an LLM gateway . Defaults to <code>https://aws-external-anthropic.{AWS_REGION}.api.aws</code>

Variable	Purpose
ANTHROPIC_AWS_WORKSPACE_ID	Required for Claude Platform on AWS . Sent on every request as the <code>anthropic-workspace-id</code> header
ANTHROPIC_BASE_URL	Override the API endpoint to route requests through a proxy or gateway. When set to a non-first-party host, MCP tool search is disabled by default. Set <code>ENABLE_TOOL_SEARCH=true</code> if your proxy forwards <code>tool_reference</code> blocks
ANTHROPIC_BEDROCK_BASE_URL	Override the Bedrock endpoint URL. Use for custom Bedrock endpoints or when routing through an LLM gateway . See Amazon Bedrock
ANTHROPIC_BEDROCK_MANTLE_BASE_URL	Override the Bedrock Mantle endpoint URL. See Mantle endpoint
ANTHROPIC_BEDROCK_SERVICE_TIER	Bedrock service tier (<code>default</code> , <code>flex</code> , or <code>priority</code>). Sent as the <code>X-Amzn-Bedrock-Service-Tier</code> header. See Amazon Bedrock
ANTHROPIC_BETAS	Comma-separated list of additional <code>anthropic-beta</code> header values to include in API requests. Claude Code already sends the beta headers it needs; use this to opt into an Anthropic API beta before Claude Code adds native support. Unlike the <code>--betas</code> flag , which requires API key authentication, this variable works with all auth methods including Claude.ai subscription
ANTHROPIC_CUSTOM_HEADERS	Custom headers to add to requests (<code>Name: Value</code> format, newline-separated for multiple headers)
ANTHROPIC_CUSTOM_MODEL_OPTION	Model ID to add as a custom entry in the <code>/model</code> picker. Use this to make a non-standard or gateway-specific model selectable without replacing built-in aliases. See Model configuration
ANTHROPIC_CUSTOM_MODEL_OPTION_DESCRIPTION	Display description for the custom model entry in the <code>/model</code> picker. Defaults to <code>Custom model (<model-id>)</code> when not set
ANTHROPIC_CUSTOM_MODEL_OPTION_NAME	Display name for the custom model entry in the <code>/model</code> picker. Defaults to the model ID when not set
ANTHROPIC_CUSTOM_MODEL_OPTION_SUPPORTED_CAPABILITIES	See Model configuration
ANTHROPIC_DEFAULT_FABLE_MODEL	See Model configuration
ANTHROPIC_DEFAULT_FABLE_MODEL_DESCRIPTION	See Model configuration
ANTHROPIC_DEFAULT_FABLE_MODEL_NAME	See Model configuration
ANTHROPIC_DEFAULT_FABLE_MODEL_SUPPORTED_CAPABILITIES	See Model configuration

Variable	Purpose
ANTHROPIC_DEFAULT_HAIKU_MODEL	See Model configuration
ANTHROPIC_DEFAULT_HAIKU_MODEL_DESCRIPTION	See Model configuration
ANTHROPIC_DEFAULT_HAIKU_MODEL_NAME	See Model configuration
ANTHROPIC_DEFAULT_HAIKU_SUPPORTED_CAPABILITIES	See Model configuration
ANTHROPIC_DEFAULT_OPENAI_MODEL	See Model configuration
ANTHROPIC_DEFAULT_OPENAI_MODEL_DESCRIPTION	See Model configuration
ANTHROPIC_DEFAULT_OPENAI_MODEL_NAME	See Model configuration
ANTHROPIC_DEFAULT_OPENAI_SUPPORTED_CAPABILITIES	See Model configuration
ANTHROPIC_DEFAULT_OPENAI_MODEL_NAME	See Model configuration
ANTHROPIC_DEFAULT_OPENAI_MODEL_DESCRIPTION	See Model configuration
ANTHROPIC_DEFAULT_OPENAI_MODEL_NAME	See Model configuration
ANTHROPIC_DEFAULT_OPENAI_SUPPORTED_CAPABILITIES	See Model configuration
ANTHROPIC_DEFAULT_OPENAI_MODEL_NAME	See Model configuration
ANTHROPIC_DEFAULT_OPENAI_MODEL_DESCRIPTION	See Model configuration
ANTHROPIC_DEFAULT_OPENAI_MODEL_NAME	See Model configuration
ANTHROPIC_DEFAULT_OPENAI_SUPPORTED_CAPABILITIES	See Model configuration
ANTHROPIC_DEFAULT_OPENAI_MODEL_NAME	See Model configuration
ANTHROPIC_DEFAULT_OPENAI_MODEL_DESCRIPTION	See Model configuration
ANTHROPIC_DEFAULT_OPENAI_MODEL_NAME	See Model configuration
ANTHROPIC_DEFAULT_OPENAI_SUPPORTED_CAPABILITIES	See Model configuration
ANTHROPIC_FOUNDRY_API_KEY	API key for Microsoft Foundry authentication (see Microsoft Foundry)
ANTHROPIC_FOUNDRY_BASE_URL	Full base URL for the Foundry resource (for example, <code>https://my-resource.services.ai.azure.com/anthropic</code>). Alternative to <code>ANTHROPIC_FOUNDRY_RESOURCE</code> (see Microsoft Foundry)
ANTHROPIC_FOUNDRY_RESOURCE	Foundry resource name (for example, <code>my-resource</code>). Required if <code>ANTHROPIC_FOUNDRY_BASE_URL</code> is not set (see Microsoft Foundry)
ANTHROPIC_MODEL	Name of the model setting to use (see Model Configuration)
ANTHROPIC_SMALL_FAST_MODEL	[DEPRECATED] Name of Haiku-class model for background tasks

Variable	Purpose
MODEL	
ANTHROPIC_SMALL_FAST_MODEL_AWS_REGION	Override AWS region for the Haiku-class model when using Bedrock or Bedrock Mantle. On Bedrock, this only takes effect when ANTHROPIC_DEFAULT_HAIKU_MODEL or the deprecated ANTHROPIC_SMALL_FAST_MODEL is also set, since Bedrock otherwise uses the primary model for background tasks
ANTHROPIC_VERTEX_BASE_URL	Override the Vertex AI endpoint URL. Use for custom Vertex endpoints or when routing through an LLM gateway . See Google Vertex AI
ANTHROPIC_VERTEX_PROJECT_ID	GCP project ID for Vertex AI requests. Overridden by G_CLOUD_PROJECT , GOOGLE_CLOUD_PROJECT , or the project in your GOOGLE_APPLICATION_CREDENTIALS credential file. See Google Vertex AI
ANTHROPIC_WORKSPACE_ID	Workspace ID for workload identity federation . Set this when your federation rule is scoped to more than one workspace so the token exchange knows which workspace to target
API_FORCE_IDLE_TIMEOUT	Override the 5-minute idle timeout that aborts a streaming model response when no bytes arrive. Set to 0 to disable the timeout, for example when a slow gateway or local model pauses longer than 5 minutes between chunks. Set to 1 to keep the timeout on every provider. When unset, the timeout is inactive on direct Anthropic API and Claude Platform on AWS connections, where Claude Code's own byte-level stream watchdog runs, and active on every other provider, including Vertex AI , Foundry , Mantle , Bedrock , and gateway connections, so a stalled stream aborts instead of hanging. As of v2.1.169
API_TIMEOUT_MS	Timeout for API requests in milliseconds (default: 600000, or 10 minutes; maximum: 2147483647). Increase this when requests time out on slow networks or when routing through a proxy. Values above the maximum overflow the underlying timer and cause requests to fail immediately
AWS_BEARER_TOKEN_BEDROCK	Bedrock API key for authentication (see Bedrock API keys)
BASH_DEFAULT_TIMEOUT_MS	Default timeout for long-running bash commands (default: 120000, or 2 minutes)
BASH_MAX_OUTPUT_LENGTH	Maximum number of characters in bash outputs before the full output is saved to a file and Claude receives the path plus a short preview. See Bash tool behavior
BASH_MAX_TIMEOUT_MS	Maximum timeout the model can set for long-running bash commands (default: 600000, or 10 minutes)
CCR_FORCE_BUNDLE	Set to 1 to force claude --remote to bundle and upload your local repository even when GitHub access is available
CLAUDECODE	Set to 1 in subprocesses Claude Code spawns (Bash and PowerShell tools, tmux sessions, hook commands, status line commands, stdio MCP server subprocesses). IDE extensions also set this in their integrated terminals. Use to detect when a script is running inside a subprocess spawned by Claude Code. To check whether the current process was spawned directly by a tool call or hook, rather than inside a stdio MCP server that Claude Code started, use CLAUDE_CODE_CHILD_SESSION instead
CLAUDE_AGENT_SDK_DISABLE	Set to 1 to disable all built-in subagent types such as Explore and Plan. Only applies in

Variable	Purpose
<code>BLE_BUILTIN_AGENTS</code>	non-interactive mode (the <code>-p</code> flag). Useful for SDK users who want a blank slate
<code>CLAUDE_AGENT_SDK_MCP_NO_PREFIX</code>	Set to <code>1</code> to skip the <code>mcp__<server>__</code> prefix on tool names from SDK-created MCP servers. Tools use their original names. SDK usage only
<code>CLAUDE_ASYNC_AGENT_STALL_TIMEOUT_MS</code>	Stall timeout in milliseconds for background subagents. Default <code>600000</code> (10 minutes). The timer resets on each streaming progress event; if no progress arrives within the window, the subagent is aborted and the task is marked failed, surfacing any partial result to the parent
<code>CLAUDE_AUTOCOMPACT_PERCENT_OVERRIDE</code>	Set the percentage (1-100) of the auto-compaction window at which auto-compaction triggers. Use lower values like <code>50</code> to compact earlier. This variable only causes earlier compaction when Claude Code compacts proactively: when <code>CLAUDE_CODE_AUTO_COMPACT_WINDOW</code> is set, in cloud sessions , and on Sonnet 4.6 and Opus 4.6 without extended context , which compact at the 200K boundary by default. In other cases, such as a local session on Opus 4.8 or any model with extended context, auto-compaction triggers when the conversation reaches the model's context limit. The override can only lower the threshold, so values above the default have no effect. Applies to both main conversations and subagents
<code>CLAUDE_AUTO_BACKGROUND_TASKS</code>	Set to <code>1</code> to force-enable automatic backgrounding of long-running agent tasks. When enabled, subagents are moved to the background after running for approximately two minutes
<code>CLAUDE_AX_SCREEN_READER</code>	Set to <code>1</code> to render screen-reader friendly output: flat text without decorative borders or animations. Set to <code>0</code> to force screen-reader mode off even when <code>axScreenReader</code> is <code>true</code> . The <code>--ax-screen-reader</code> flag takes precedence. Requires Claude Code v2.1.181 or later
<code>CLAUDE_BASH_MAINTAIN_PROJECT_WORKING_DIR</code>	Return to the original working directory after each Bash or PowerShell command in the main session
<code>CLAUDE_CLIENT_PRESENCE_FILE</code>	Path to a file that an external tool, such as a screen-lock listener, creates when you unlock your screen and deletes when you lock it. While the file exists, Claude Code skips Remote Control mobile push notifications , so you stop getting pushes while you are actively using the computer. When the file is absent or unreadable, notifications are sent as normal. Claude Code checks the file once per push-triggering event rather than polling it. Requires Claude Code v2.1.181 or later
<code>CLAUDE_CODE_ACCESSIBILITY</code>	Set to <code>1</code> to keep the native terminal cursor visible and disable the inverted-text cursor indicator. Allows screen magnifiers like macOS Zoom to track cursor position
<code>CLAUDE_CODE_ADDITIONAL_DIRECTORIES_CLAUDE_MD</code>	Set to <code>1</code> to load memory files from directories specified with <code>--add-dir</code> . Loads <code>CLAUDE.md</code> , <code>.claude/CLAUDE.md</code> , <code>.claude/rules/*.md</code> , and <code>CLAUDE.local.md</code> . By default, additional directories do not load memory files
<code>CLAUDE_CODE_ALT_SCREEN_FULL_REPAINT</code>	Set to <code>1</code> to repaint the entire screen on every frame in fullscreen rendering instead of sending incremental updates. Use this if fullscreen mode shows stale or misplaced text fragments. Claude Code enables this automatically for background sessions and agent view on Windows
<code>CLAUDE_CODE_ALWAYS_ENABLE EFFORT</code>	Set to <code>1</code> to send the effort parameter with every request, even when Claude Code does not recognize the model ID as effort-capable. Use this when routing through an LLM gateway or third-party provider that serves models under custom identifiers.

Variable	Purpose
	Models that reject the effort parameter at the API, including Claude 3 models, Sonnet 4.0 and 4.5, Opus 4.0 and 4.1, and Haiku 4.5, are still excluded so requests do not fail
CLAUDE_CODE_API_KEY_HELPER_TTL_MS	Interval in milliseconds at which credentials should be refreshed (when using <code>apiKeyHelper</code>)
CLAUDE_CODE_ARTIFACT_AUTO_OPEN	Set to <code>0</code> to stop Claude Code from opening the browser automatically when a new artifact is published. Republishing an existing artifact does not open the browser regardless of this setting
CLAUDE_CODE_ATTRIBUTION_HEADER	Set to <code>0</code> to omit the attribution block (client version and prompt fingerprint) from the start of the system prompt. Disabling it improves prompt-cache hit rates when routing through an LLM gateway . Anthropic API caching is unaffected
CLAUDE_CODE_AUTO_COMPACT_WINDOW	Set the context capacity in tokens used for auto-compaction calculations. Defaults to the model's context window: 200K for standard models or 1M for extended context models. Use a lower value like <code>500000</code> on a 1M model to treat the window as 500K for compaction purposes. The value is capped at the model's actual context window. <code>CLAUDE_AUTOCOMPACT_PCT_OVERRIDE</code> is applied as a percentage of this value. Setting this variable decouples the compaction threshold from the status line's <code>used_percentage</code> , which always uses the model's full context window
CLAUDE_CODE_AUTO_CONNECT_IDE	Override automatic IDE connection . By default, Claude Code connects automatically when launched inside a supported IDE's integrated terminal. Set to <code>false</code> to prevent this. Set to <code>true</code> to force a connection attempt when auto-detection fails, such as when tmux obscures the parent terminal. Takes precedence over the <code>autoConnectIde</code> global config setting
CLAUDE_CODE_CERT_STORE	Comma-separated list of CA certificate sources for TLS connections. <code>bundled</code> is the Mozilla CA set shipped with Claude Code. <code>system</code> is the operating system trust store. Default is <code>bundled,system</code>
CLAUDE_CODE_CHILD_SESSION	Set to <code>1</code> in subprocesses Claude Code spawns via the Bash, PowerShell, and Monitor tools, hook commands, and status line commands. Not set for stdio MCP server subprocesses, which are long-lived and outlive the session that spawned them. Unlike <code>CLAUDECODE</code> , this is only set by Claude Code itself when it launches a subprocess and not by IDE extensions, so it reliably distinguishes a nested session from a top-level <code>claude</code> launched in an IDE-integrated terminal. A nested interactive <code>claude</code> TUI started this way is automatically excluded from <code>--resume</code> , <code>--continue</code> , up-arrow history, and the <code>claude agents</code> list. Non-interactive <code>claude -p</code> sessions still persist. Set <code>CLAUDE_CODE_FORCE_SESSION_PERSISTENCE=1</code> to override this exclusion. Requires Claude Code v2.1.172 or later
CLAUDE_CODE_CLIENT_CERT	Path to client certificate file for mTLS authentication
CLAUDE_CODE_CLIENT_KEY	Path to client private key file for mTLS authentication
CLAUDE_CODE_CLIENT_KEY_PASSPHRASE	Passphrase for encrypted CLAUDE_CODE_CLIENT_KEY (optional)
CLAUDE_CODE_CONNECT_TIMEOUT_MS	Removed in v2.1.186 and now a no-op. Previously set a separate timeout for the connect, TLS, and response-header phase of a streaming API request. Use <code>API_TIMEOUT_MS</code> for

Variable	Purpose
	the per-request timeout
CLAUDE_CODE_DEBUG_LOG S_DIR	Override the debug log file path. Despite the name, this is a file path, not a directory. Requires debug mode to be enabled separately via <code>--debug</code> , <code>/debug</code> , or the <code>DEBUG</code> environment variable: setting this variable alone does not enable logging. The <code>--debug-file</code> flag does both at once. Defaults to <code>~/.claude/debug/<session-id>.txt</code>
CLAUDE_CODE_DEBUG_LOG _LEVEL	Minimum log level written to the debug log file. Values: <code>verbose</code> , <code>debug</code> (default), <code>info</code> , <code>warn</code> , <code>error</code> . Set to <code>verbose</code> to include high-volume diagnostics like full status line command output, or raise to <code>error</code> to reduce noise
CLAUDE_CODE_DISABLE_1 M_CONTEXT	Set to <code>1</code> to disable 1M context window support. When set, 1M model variants are unavailable in the model picker. Useful for enterprise environments with compliance requirements
CLAUDE_CODE_DISABLE_A DAPTIVE_THINKING	Set to <code>1</code> to disable adaptive reasoning on Opus 4.6 and Sonnet 4.6 and fall back to the fixed thinking budget controlled by <code>MAX_THINKING_TOKENS</code> . From v2.1.111, has no effect on Fable 5, or on Opus 4.7 and later, which always use adaptive reasoning
CLAUDE_CODE_DISABLE_A DVISOR_TOOL	Set to <code>1</code> to disable the advisor tool . The <code>/advisor</code> command and <code>--advisor</code> flag become unavailable and any configured <code>advisorModel</code> is ignored. Requires Claude Code v2.1.98 or later
CLAUDE_CODE_DISABLE_A GENT_VIEW	Set to <code>1</code> to turn off background agents and agent view : <code>claude agents</code> , <code>--bg</code> , <code>/background</code> , and the on-demand supervisor. Equivalent to the <code>disableAgentView</code> setting
CLAUDE_CODE_DISABLE_A LTERNATE_SCREEN	Set to <code>1</code> to disable fullscreen rendering and use the classic main-screen renderer. The conversation stays in your terminal's native scrollback so <code>Cmd+f</code> and tmux copy mode work as usual. Takes precedence over <code>CLAUDE_CODE_NO_FLICKER</code> and the <code>tui</code> setting. You can also switch with <code>/tui default</code> . Does not apply to background sessions opened from agent view , which always use fullscreen rendering
CLAUDE_CODE_DISABLE_A RTIFACT	Set to <code>1</code> to disable the Artifact tool, which publishes session output as a private web page on claude.ai. Equivalent to the <code>disableArtifact</code> setting
CLAUDE_CODE_DISABLE_A TTACHMENTS	Set to <code>1</code> to disable attachment processing. File mentions with <code>@</code> syntax are sent as plain text instead of being expanded into file content
CLAUDE_CODE_DISABLE_A UTO_MEMORY	Set to <code>1</code> to disable auto memory . Set to <code>0</code> to force auto memory on even when <code>--bare mode</code> or <code>autoMemoryEnabled: false</code> would otherwise disable it. When disabled, Claude does not create or load auto memory files
CLAUDE_CODE_DISABLE_B ACKGROUND_TASKS	Set to <code>1</code> to disable all background task functionality, including the <code>run_in_background</code> parameter on Bash and subagent tools, auto-backgrounding, and the Ctrl+B shortcut
CLAUDE_CODE_DISABLE_B UNDLED_SKILLS	Set to <code>1</code> to disable the skills and workflows that ship with Claude Code: bundled skills and workflows are removed entirely, while built-in slash commands like <code>/init</code> stay typable but are hidden from the model. Skills from plugins, <code>.claude/skills/</code> , and <code>.claude/commands/</code> are unaffected. Equivalent to the <code>disableBundledSkills</code> setting; <code>0</code> does not override it
CLAUDE_CODE_DISABLE_C	Set to <code>1</code> to prevent loading any CLAUDE.md memory files into context, including user,

Variable	Purpose
LAUDE_MDS	project, and auto-memory files
CLAUDE_CODE_DISABLE_CRON	Set to <code>1</code> to disable scheduled tasks . The <code>/loop</code> skill and cron tools become unavailable and any already-scheduled tasks stop firing, including tasks that are already running mid-session
CLAUDE_CODE_DISABLE_EXPERIMENTAL_BETAS	Set to <code>1</code> to strip Anthropic-specific <code>anthropic-beta</code> request headers and beta tool-schema fields (such as <code>defer_loading</code> and <code>eager_input_streaming</code>) from API requests. Use this when a proxy gateway rejects requests with errors like “Unexpected value(s) for the <code>anthropic-beta</code> header” or “Extra inputs are not permitted”. Standard fields (<code>name</code> , <code>description</code> , <code>input_schema</code> , <code>cache_control</code>) are preserved
CLAUDE_CODE_DISABLE_FAST_MODE	Set to <code>1</code> to disable <u>fast mode</u>
CLAUDE_CODE_DISABLE_FEEDBACK_SURVEY	Set to <code>1</code> to disable the “How is Claude doing?” session quality surveys. Surveys are also disabled when <code>DISABLE_TELEMETRY</code> , <code>DO_NOT_TRACK</code> , or <code>CLAUDE_CODE_DISABLE_NONESSENTIAL_TRAFFIC</code> is set, unless <code>CLAUDE_CODE_ENABLE_FEEDBACK_SURVEY_FOR_OTEL</code> opts back in. To set a sample rate instead of disabling outright, use the <code>feedbackSurveyRate</code> setting. See <u>Session quality surveys</u>
CLAUDE_CODE_DISABLE_FILE_CHECKPOINTING	Set to <code>1</code> to disable file <u>checkpointing</u> . The <code>/rewind</code> command will not be able to restore code changes
CLAUDE_CODE_DISABLE_GIT_INSTRUCTIONS	Set to <code>1</code> to remove built-in commit and PR workflow instructions and the git status snapshot from Claude’s system prompt. Useful when using your own git workflow skills. Takes precedence over the <code>includeGitInstructions</code> setting when set
CLAUDE_CODE_DISABLE_LEGACY_MODEL_REMAP	Set to <code>1</code> to prevent automatic remapping of Opus 4.0 and 4.1 to the current Opus version on the Anthropic API. Use when you intentionally want to pin an older model. The remap does not run on Bedrock, Vertex, or Foundry
CLAUDE_CODE_DISABLE_MOUSE	Set to <code>1</code> to disable mouse tracking in <u>fullscreen rendering</u> . Keyboard scrolling with <code>PgUp</code> and <code>PgDn</code> still works. Use this to keep your terminal’s native copy-on-select behavior
CLAUDE_CODE_DISABLE_NONESSENTIAL_TRAFFIC	Equivalent of setting <code>DISABLE_AUTOUPDATER</code> , <code>DISABLE_FEEDBACK_COMMAND</code> , <code>DISABLE_ERROR_REPORTING</code> , and <code>DISABLE_TELEMETRY</code>
CLAUDE_CODE_DISABLE_NONSTREAMING_FALLBACK	Set to <code>1</code> to disable the non-streaming fallback when a streaming request fails mid-stream. Streaming errors propagate to the retry layer instead. Useful when a proxy or gateway causes the fallback to produce duplicate tool execution
CLAUDE_CODE_DISABLE_OFFICIAL_MARKETPLACE_AUTOINSTALL	Set to <code>1</code> to skip automatic addition of the official plugin marketplace on first run
CLAUDE_CODE_DISABLE_POLICY_SKILLS	Set to <code>1</code> to skip loading skills from the system-wide managed skills directory. Useful for container or CI sessions that should not load operator-provisioned skills
CLAUDE_CODE_DISABLE_TERMINAL_TITLE	Set to <code>1</code> to disable automatic terminal title updates based on conversation context. In Agent SDK and <code>claude -p</code> sessions, this also skips the background Haiku request that generates the session title

Variable	Purpose
CLAUDE_CODE_DISABLE_THINKING	Set to <code>1</code> to omit the <code>thinking</code> parameter from API requests entirely. This is a compatibility option for proxies and gateways that reject the parameter. The variable's behavior is unchanged from earlier versions; on models that think by default, omitting the parameter means the model may still think. To explicitly disable extended thinking on the Anthropic API, use <code>MAX_THINKING_TOKENS=0</code> instead, which is also ineffective on Fable 5 since it cannot have thinking turned off. On third-party providers , <code>0</code> likewise omits the parameter, so the two variables behave the same there
CLAUDE_CODE_DISABLE_VIRTUAL_SCROLL	Set to <code>1</code> to disable virtual scrolling in fullscreen rendering and render every message in the transcript. Use this if scrolling in fullscreen mode shows blank regions where messages should appear
CLAUDE_CODE_DISABLE_WORKFLOWS	Set to <code>1</code> to disable workflows . Equivalent to the <code>disableWorkflows</code> setting
CLAUDE_CODE_EFFORT_LEVEL	Set the effort level for supported models. Values: <code>low</code> , <code>medium</code> , <code>high</code> , <code>xhigh</code> , <code>max</code> , or <code>auto</code> to use the model default. Available levels depend on the model. Takes precedence over <code>/effort</code> and the <code>effortLevel</code> setting. See Adjust effort level
CLAUDE_CODE_ENABLE_AUTO_MODE	Set to <code>1</code> to make auto mode available on Amazon Bedrock, Google Cloud Vertex AI, and Microsoft Foundry. Requires Claude Code v2.1.158 or later. Has no effect on the Anthropic API, where auto mode is available by default. See Enable auto mode on Bedrock, Vertex AI, or Foundry
CLAUDE_CODE_ENABLE_AWAY_SUMMARY	Override session recap availability. Set to <code>0</code> to force recaps off regardless of the <code>/config</code> toggle. Set to <code>1</code> to force recaps on when <code>awaySummaryEnabled</code> is <code>false</code> . Takes precedence over the setting and <code>/config</code> toggle
CLAUDE_CODE_ENABLE_BACKGROUND_PLUGIN_REFRESH	Set to <code>1</code> to refresh plugin state at turn boundaries in non-interactive mode after a background install completes. Off by default because the refresh changes the system prompt mid-session, which invalidates prompt caching for that turn
CLAUDE_CODE_ENABLE_FEEDBACK_SURVEY_FOR_OTEL	Set to <code>1</code> to route the “How is Claude doing?” session quality survey to your own OpenTelemetry collector when Anthropic-bound nonessential traffic is blocked. Survey ratings are emitted only as OTEL events to your configured collector. No survey data is sent to Anthropic in this mode. Applies when <code>CLAUDE_CODE_DISABLE_NONESSENTIAL_TRAFFIC</code> , <code>DISABLE_TELEMETRY</code> , or <code>DO_NOT_TRACK</code> is set, and has no effect otherwise. <code>CLAUDE_CODE_DISABLE_FEEDBACK_SURVEY</code> and the organization product feedback policy take precedence
CLAUDE_CODE_ENABLE_FINE_GRAINED_TOOL_STREAMING	Controls whether tool call inputs stream from the API as Claude generates them. With this off, a large tool input such as a long file write arrives only after Claude finishes generating it, which can look like it's hanging. Enabled by default on the Anthropic API. On Bedrock and Vertex, enabled per model where the deployed container supports it. Set to <code>0</code> to opt out. Set to <code>1</code> to force on when routing through a proxy via <code>ANTHROPIC_BASE_URL</code> , <code>ANTHROPIC_VERTEX_BASE_URL</code> , or <code>ANTHROPIC_BEDROCK_BASE_URL</code> . Off by default on Foundry and gateway connections
CLAUDE_CODE_ENABLE_GATEWAY_MODEL_DISCOVERY	Set to <code>1</code> to populate the <code>/model</code> picker from your gateway's <code>/v1/models</code> endpoint when <code>ANTHROPIC_BASE_URL</code> points at an Anthropic-compatible gateway such as LiteLLM, Kong, or an internal proxy. Off by default because gateways backed by a shared API key would otherwise show every user every model the key can access. Discovered models are still filtered by an <code>availableModels</code> allowlist the session receives; deliver

Variable	Purpose
	the list through MDM or a managed settings file , since server-managed delivery is not available on gateway configurations
CLAUDE_CODE_ENABLE_OPUS_4_7_FAST_MODE	Removed in v2.1.142, when the fast mode default moved from Opus 4.6 to Opus 4.7
CLAUDE_CODE_ENABLE_PROMPT_SUGGESTION	Set to <code>false</code> to disable prompt suggestions (the “Prompt suggestions” toggle in <code>/config</code>). These are the grayed-out predictions that appear in your prompt input after Claude responds. See <u>Prompt suggestions</u>
CLAUDE_CODE_ENABLE_TASKS	Controls whether sessions use the structured Task tools (<code>TaskCreate</code> , <code>TaskUpdate</code> , <code>TaskGet</code> , <code>TaskList</code>) or the legacy <code>TodoWrite</code> tool. As of Claude Code v2.1.142, Task tools are the default in all modes. Set to <code>0</code> to revert to <code>TodoWrite</code> . See <u>Task list</u> and <u>Migrate to Task tools</u>
CLAUDE_CODE_ENABLE_TELEMETRY	Set to <code>1</code> to enable OpenTelemetry data collection for metrics and logging. Required before configuring OTel exporters. See <u>Monitoring</u>
CLAUDE_CODE_EXIT_AFTER_STOP_DELAY	Time in milliseconds to wait after the query loop becomes idle before automatically exiting. Useful for automated workflows and scripts using SDK mode
CLAUDE_CODE_EXPERIMENTAL_AGENT_TEAMS	Set to <code>1</code> to enable <u>agent teams</u> . Agent teams are experimental and disabled by default
CLAUDE_CODE_EXTRA_BODY	JSON object to merge into the top level of every API request body. Useful for passing provider-specific parameters that Claude Code does not expose directly
CLAUDE_CODE_FILE_READ_MAX_OUTPUT_TOKENS	Override the default token limit for file reads. Useful when you need to read larger files in full
CLAUDE_CODE_FORCE_SESSION_PERSISTENCE	Set to <code>1</code> to force transcript persistence, prompt history, and <code>claude agents</code> registration even when this <code>claude</code> was launched from inside another Claude Code session. Use when an inherited <code>CLAUDE_CODE_CHILD_SESSION</code> value, for example from a <code>screen</code> session or a background launcher first started by Claude Code's Bash tool, causes a genuine top-level session to be misclassified as nested. As of v2.1.178, Claude Code detects the tmux case automatically and ignores the inherited marker, so tmux no longer needs this variable. Also honored on v2.1.169 and earlier; has no effect on v2.1.170 and v2.1.171, where the nested-session detection it overrides was removed
CLAUDE_CODE_FORCE_STRIKETHROUGH	Set to <code>1</code> to force strikethrough rendering for <code>~~text~~</code> in Claude's responses when your terminal supports it but is not auto-detected, such as over SSH without <code>TERM_PROGRAM</code> forwarded. Without this, undetected terminals show the literal <code>~~</code> markers instead of rendering the text as strikethrough. Requires Claude Code v2.1.186 or later
CLAUDE_CODE_FORCE_SYNC_OUTPUT	Set to <code>1</code> to force-enable DEC private mode 2026 <u>synchronized output</u> when your terminal supports it but is not auto-detected. Useful for emulators such as Emacs <code>eat</code> that implement BSU/ESU but do not reply to the capability probe. Has no effect under tmux
CLAUDE_CODE_FORK_SUBAGENT	Set to <code>1</code> to let Claude spawn <u>forked subagents</u> , or <code>0</code> to disable them, overriding any server-side rollout. When enabled, Claude can request the <code>fork</code> subagent type to spawn a fork, a subagent that inherits the full conversation context instead of starting fresh. Spawns without a subagent type still use the general-purpose subagent, and all

Variable	Purpose
	subagent spawns run in the background. The explicit <code>/fork</code> command works without this variable. Works in interactive mode and via the SDK or <code>claude -p</code>
<code>CLAUDE_CODE_GIT_BASH_PATH</code>	Windows only: path to the Git Bash executable (<code>bash.exe</code>). Use when Git Bash is installed but not in your PATH. See Windows setup
<code>CLAUDE_CODE_GLOB_HIDDEN</code>	Set to <code>false</code> to exclude dotfiles from results when Claude invokes the Glob tool . Included by default. Does not affect <code>@</code> file autocomplete, <code>ls</code> , Grep, or Read
<code>CLAUDE_CODE_GLOB_NO_IGNORE</code>	Set to <code>false</code> to make the Glob tool respect <code>.gitignore</code> patterns. By default, Glob returns all matching files including gitignored ones. Does not affect <code>@</code> file autocomplete, which has its own respectGitignore setting
<code>CLAUDE_CODE_GLOB_TIMEOUT_SECONDS</code>	Timeout in seconds for Glob tool file discovery. Defaults to 20 seconds on most platforms and 60 seconds on WSL
<code>CLAUDE_CODE_HIDE_CWD</code>	Set to <code>1</code> to hide the working directory in the startup logo. Useful for screenshares or recordings where the path exposes your OS username
<code>CLAUDE_CODE_IDE_HOST_OVERRIDE</code>	Override the host address used to connect to the IDE extension. By default Claude Code auto-detects the correct address, including WSL-to-Windows routing
<code>CLAUDE_CODE_IDE_SKIP_AUTO_INSTALL</code>	Skip auto-installation of IDE extensions. Equivalent to setting autoInstallIdeExtension to <code>false</code>
<code>CLAUDE_CODE_IDE_SKIP_VALID_CHECK</code>	Set to <code>1</code> to skip validation of IDE lockfile entries during connection. Use when auto-connect fails to find your IDE despite it running
<code>CLAUDE_CODE_MAX_CONTEXT_TOKENS</code>	Override the context window size Claude Code assumes for the active model. Only takes effect when <code>DISABLE_COMPACT</code> is also set. Use this when routing to a model through <code>ANTHROPIC_BASE_URL</code> whose context window does not match the built-in size for its name
<code>CLAUDE_CODE_MAX_OUTPUT_TOKENS</code>	Set the maximum number of output tokens for most requests. Defaults and caps vary by model; see max output tokens . Increasing this value reduces the effective context window available before auto-compaction triggers
<code>CLAUDE_CODE_MAX_RETRIES</code>	Override the number of times to retry failed API requests (default: 10). Capped at 15 as of v2.1.186. For unattended sessions that need to wait through longer outages, set <code>CLAUDE_CODE_RETRY_WATCHDOG</code> instead
<code>CLAUDE_CODE_MAX_TOOL_USE_CONCURRENCY</code>	Maximum number of read-only tools and subagents that can execute in parallel (default: 10). Higher values increase parallelism but consume more resources
<code>CLAUDE_CODE_MAX_TURNS</code>	Cap the number of agentic turns when no explicit limit is passed. Equivalent to passing --max-turns , which takes precedence when both are set. A value that is not a positive integer is rejected at startup with an error rather than treated as no cap
<code>CLAUDE_CODE_MCP_ALLOW_LIST_ENV</code>	Set to <code>1</code> to spawn stdio MCP servers with only a safe baseline environment plus the server's configured <code>env</code> , instead of inheriting your shell environment
<code>CLAUDE_CODE_MCP_TOOL_IDLE_TIMEOUT</code>	Idle timeout in milliseconds for remote MCP tool calls (default: 300000, or 5 minutes). When an HTTP, SSE, WebSocket, or claude.ai connector MCP server sends no response and no progress notification for this long, the tool call aborts with an error instead of waiting for the overall <code>MCP_TOOL_TIMEOUT</code> . Set to <code>0</code> to disable the idle check. Values

Variable	Purpose
	below 1000 are raised to one second, and the value is capped at the effective <code>MCP_T00L_TIMEOUT</code> . Does not apply to stdio or IDE servers. Requires Claude Code v2.1.187 or later
<code>CLAUDE_CODE_NATIVE_CURSOR</code>	Set to <code>1</code> to show the terminal's own cursor at the input caret instead of a drawn block. The cursor respects the terminal's blink, shape, and focus settings
<code>CLAUDE_CODE_NEW_INIT</code>	Set to <code>1</code> to make <code>/init</code> run an interactive setup flow. The flow asks which files to generate, including CLAUDE.md, skills, and hooks, before exploring the codebase and writing them. Without this variable, <code>/init</code> generates a CLAUDE.md automatically without prompting
<code>CLAUDE_CODE_NO_FLICKER</code>	Set to <code>1</code> to enable fullscreen rendering , a research preview that reduces flicker and keeps memory flat in long conversations. Equivalent to the <code>tui</code> setting; you can also switch with <code>/tui fullscreen</code>
<code>CLAUDE_CODE_OAUTH_REFRESH_TOKEN</code>	OAuth refresh token for Claude.ai authentication. When set, <code>claude auth login</code> exchanges this token directly instead of opening a browser. Requires <code>CLAUDE_CODE_OAUTH_SCOPES</code> . Useful for provisioning authentication in automated environments
<code>CLAUDE_CODE_OAUTH_SCOPES</code>	Space-separated OAuth scopes the refresh token was issued with, such as <code>"user:profile user:inference user:sessions:claude_code"</code> . Required when <code>CLAUDE_CODE_OAUTH_REFRESH_TOKEN</code> is set
<code>CLAUDE_CODE_OAUTH_TOKEN</code>	OAuth access token for Claude.ai authentication. Alternative to <code>/login</code> for SDK and automated environments. Takes precedence over keychain-stored credentials. Generate one with <code>claude setup-token</code>
<code>CLAUDE_CODE_OPUS_4_6_FAST_MODE_OVERRIDE</code>	Removed in v2.1.160 and now a no-op. Previously pinned fast mode to Claude Opus 4.6 instead of the current default
<code>CLAUDE_CODE_OTEL_DIAG_STDERR</code>	Set to <code>1</code> to write OpenTelemetry exporter diagnostic errors to stderr. By default these errors only appear with <code>--debug</code> , so a misconfigured exporter such as a Prometheus port collision otherwise fails silently. Requires Claude Code v2.1.179 or later. See <u>Monitoring</u>
<code>CLAUDE_CODE_OTEL_FLUSH_TIMEOUT_MS</code>	Timeout in milliseconds for flushing pending OpenTelemetry spans (default: 5000). See <u>Monitoring</u>
<code>CLAUDE_CODE_OTEL_HEADERS_HELPER_DEBOUNCE_MS</code>	Interval for refreshing dynamic OpenTelemetry headers in milliseconds (default: 1740000 / 29 minutes). See <u>Dynamic headers</u>
<code>CLAUDE_CODE_OTEL_SHUTDOWN_TIMEOUT_MS</code>	Timeout in milliseconds for the OpenTelemetry exporter to finish on shutdown (default: 2000). Increase if metrics are dropped at exit. See <u>Monitoring</u>
<code>CLAUDE_CODE_PACKAGE_MANAGER_AUTO_UPDATE</code>	Set to <code>1</code> to let Claude Code run your package manager's upgrade command in the background when a new version is available. Applies to Homebrew and WinGet installations. Other package managers continue to show the upgrade command without running it. See <u>Auto updates</u>
<code>CLAUDE_CODE_PERFORCE_MODE</code>	Set to <code>1</code> to enable Perforce-aware write protection. When set, Edit, Write, and NotebookEdit fail with a <code>p4 edit <file></code> hint if the target file lacks the owner-write

Variable	Purpose
	bit, which Perforce clears on synced files until <code>p4 edit</code> opens them. This prevents Claude Code from bypassing Perforce change tracking
<code>CLAUDE_CODE_PLUGIN_CACHE_DIR</code>	Override the plugins root directory. Despite the name, this sets the parent directory, not the cache itself: marketplaces and the plugin cache live in subdirectories under this path. Defaults to <code>~/.claude/plugins</code>
<code>CLAUDE_CODE_PLUGIN_GIT_TIMEOUT_MS</code>	Timeout in milliseconds for git operations when installing or updating plugins (default: 120000). Increase this value for large repositories or slow network connections. See Git operations time out
<code>CLAUDE_CODE_PLUGIN_KEEP_MARKETPLACE_ON_FAILURE</code>	Set to <code>1</code> to keep the existing marketplace cache when a <code>git pull</code> fails instead of wiping and re-cloning. Useful in offline or airgapped environments where re-cloning would fail the same way. See Marketplace updates fail in offline environments
<code>CLAUDE_CODE_PLUGIN_PREFER_HTTPS</code>	Set to <code>1</code> to clone GitHub <code>owner/repo</code> shorthand sources over HTTPS instead of SSH. Applies to plugin install and update, and to <code>/plugin marketplace add</code> and <code>update</code> . Useful in CI runners, containers, or any environment without a configured SSH key for <code>github.com</code>
<code>CLAUDE_CODE_PLUGIN_SEED_DIR</code>	Path to one or more read-only plugin seed directories, separated by <code>:</code> on Unix or <code>;</code> on Windows. Use this to bundle a pre-populated plugins directory into a container image. Claude Code registers marketplaces from these directories at startup and uses pre-cached plugins without re-cloning. See Pre-populate plugins for containers
<code>CLAUDE_CODE_POWERSHELL_RESPECT_EXECUTION_POLICY</code>	Set to <code>1</code> to stop Claude Code from passing <code>-ExecutionPolicy Bypass</code> when spawning PowerShell for tool calls, hooks, and status line commands, and respect the machine's effective execution policy instead. By default Claude Code bypasses execution policy at process scope so <code>.ps1</code> scripts and module imports work on default-Restricted Windows installs. Process-scope bypass never overrides Group Policy <code>MachinePolicy</code> or <code>UserPolicy</code> regardless of this setting
<code>CLAUDE_CODE_PRINT_BACKGROUND_WAIT_CEILING_MS</code>	Maximum time in milliseconds that non-interactive mode with the <code>-p</code> flag waits after the final turn for background subagents and workflows whose result is part of the output. Default: <code>600000</code> , or 10 minutes. When the cap is exceeded, remaining background tasks are terminated and the process exits. Set to <code>0</code> to wait indefinitely. This cap is separate from the five-second grace period that applies to plain background shells
<code>CLAUDE_CODE_PROPAGATE_TRACEPARENT</code>	Set to <code>1</code> to propagate W3C trace context when <code>ANTHROPIC_BASE_URL</code> points at a custom proxy. Propagation covers the <code>traceparent</code> header on model and HTTP MCP requests and the <code>TRACEPARENT</code> environment variable for Bash, PowerShell, and hook subprocesses. By default, propagation is enabled only when connected directly to the Anthropic API. Added in v2.1.152. See Traces (beta)
<code>CLAUDE_CODE_PROVIDER_MANAGED_BY_HOST</code>	Set by host platforms that embed Claude Code and manage model provider routing on its behalf. When set, provider-selection, endpoint, and authentication variables such as <code>CLAUDE_CODE_USE_BEDROCK</code> , <code>ANTHROPIC_BASE_URL</code> , and <code>ANTHROPIC_API_KEY</code> in settings files are ignored so user settings cannot override the host's routing. The automatic telemetry opt-out for Bedrock, Vertex, and Foundry is also skipped, so telemetry follows the standard <code>DISABLE_TELEMETRY</code> opt-out. See Default behaviors by API provider
<code>CLAUDE_CODE_PROXY_RESOLVE</code>	Set to <code>1</code> to allow the proxy to perform DNS resolution instead of the caller. Opt-in for

Variable	Purpose
OLVES_HOSTS	environments where the proxy should handle hostname resolution
CLAUDE_CODE_REMOTE	Set automatically to <code>true</code> when Claude Code is running as a cloud session . Read this from a hook or setup script to detect whether you are in a cloud environment
CLAUDE_CODE_REMOTE_SESSION_ID	Set automatically in cloud sessions to the current session's ID. Read this to construct a link back to the session transcript. See Link output back to the session
CLAUDE_CODE_RESUME_INTERRUPTED_TURN	Set to <code>1</code> to automatically resume if the previous session ended mid-turn. Used in SDK mode so the model continues without requiring the SDK to re-send the prompt
CLAUDE_CODE_RESUME_PROMPT	Override the continuation message injected when resuming a session that ended mid-turn. Defaults to <code>Continue from where you left off.</code> . Spawn scripts for long-running agents can set this to a more directive boot message. An empty string uses the default
CLAUDE_CODE_RETRY_WATCHDOG	Set to <code>1</code> for unattended sessions such as eval harnesses, CI jobs, or remote workers. Retries <code>429</code> and <code>529</code> capacity errors indefinitely instead of failing after <code>CLAUDE_CODE_MAX_RETRIES</code> attempts. The watchdog backs off up to 5 minutes between attempts, or until the limit resets when the response carries a rate-limit reset time, so a session that hits a usage limit waits out the remaining window. Requires Claude Code v2.1.186 or later
CLAUDE_CODE_SAFE_MODE	Set to <code>1</code> to start in safe mode: CLAUDE.md, skills, plugins, hooks, MCP servers, custom commands and agents, output styles, workflows, custom themes, custom keybindings, status line and file-suggestion commands, LSP servers, and auto-memory do not load, for troubleshooting a broken configuration. Managed settings policy still applies, including policy-configured hooks, status line, and file-suggestion commands; managed plugins, managed skills, managed CLAUDE.md, and policy-configured MCP servers do not. Equivalent to passing <code>--safe-mode</code> . Directly spawned child processes inherit the variable
CLAUDE_CODE_SCRIPT_CAPS	JSON object limiting how many times specific scripts may be invoked per session when <code>CLAUDE_CODE_SUBPROCESS_ENV_SCRUB</code> is set. Keys are substrings matched against the command text; values are integer call limits. For example, <code>{"deploy.sh": 2}</code> allows <code>deploy.sh</code> to be called at most twice. Matching is substring-based so shell-expansion tricks like <code>./scripts/deploy.sh \$(evil)</code> still count against the cap. Runtime fan-out via <code>xargs</code> or <code>find -exec</code> is not detected; this is a defense-in-depth control
CLAUDE_CODE_SCROLL_SPEED	Set the mouse wheel scroll multiplier in fullscreen rendering . Accepts values from 1 to 20, and fractional values below 1 such as <code>0.5</code> to slow accelerated trackpad and wheel scrolling in terminals that already amplify wheel events. Set to <code>3</code> to match <code>vim</code> if your terminal sends one wheel event per notch without amplification. Ignored in the JetBrains IDE terminal, where Claude Code uses its own scroll handling
CLAUDE_CODE_SESSION_END_HOOKS_TIMEOUT_MS	Override the time budget in milliseconds for SessionEnd hooks. Applies to session exit, <code>/clear</code> , and switching sessions via interactive <code>/resume</code> . By default the budget is 1.5 seconds, automatically raised to the highest per-hook <code>timeout</code> configured in settings files, up to 60 seconds. Timeouts on plugin-provided hooks do not raise the budget
CLAUDE_CODE_SESSION_ID	Set automatically to the current session ID in Bash and PowerShell tool subprocesses, hook command subprocesses, and stdio MCP server subprocesses. For Bash, PowerShell, and hooks this matches the <code>session_id</code> field in the hook JSON input and is updated on <code>/clear</code> . An MCP server subprocess retains the ID it was spawned with. On <code>--resume <session-id></code> it receives the resumed ID, matching hooks and Bash. On

Variable	Purpose
	<code>--continue</code> or <code>--resume</code> without an explicit ID it may receive the initial startup ID instead. Use to correlate scripts and external tools with the Claude Code session that launched them
<code>CLAUDE_CODE_SHELL</code>	Override automatic shell detection. Useful when your login shell differs from your preferred working shell (for example, <code>bash</code> vs <code>zsh</code>)
<code>CLAUDE_CODE_SHELL_PREFIX</code>	Command prefix that wraps shell commands Claude Code spawns: Bash tool calls, hook commands, status line commands, and stdio MCP server startup commands. PowerShell hooks and exec-form hooks run without the prefix. Useful for logging or auditing. Setting a bare executable path such as <code>/path/to/logger.sh</code> runs each command as <code>/path/to/logger.sh '<command>'</code> . The wrapper receives the command line as a single shell-quoted argument in <code>\$1</code> , so the wrapper must re-evaluate <code>\$1</code> with a shell, for example <code>exec bash -c "\$1"</code> . Treating <code>\$1</code> as a bare executable path breaks stdio MCP servers that pass arguments such as <code>npx -y <package></code> . For Bash tool calls, <code>\$1</code> contains the full shell invocation Claude Code assembles, including environment setup, not only the command Claude ran
<code>CLAUDE_CODE_SIMPLE</code>	Set to <code>1</code> to run with a minimal system prompt and only the Bash, file read, and file edit tools. MCP tools from <code>--mcp-config</code> are still available. Disables auto-discovery of hooks, skills, plugins, MCP servers, auto memory, and CLAUDE.md. OAuth tokens and keychain credentials are not read, so Anthropic authentication must come from <code>ANTHROPIC_API_KEY</code> or an <code>apiKeyHelper</code> in <code>--settings</code> . Equivalent to passing <code>--bare</code>
<code>CLAUDE_CODE_SIMPLE_SYSTEM_PROMPT</code>	Set to <code>1</code> to use a shorter system prompt and abbreviated tool descriptions on any model. Set to <code>0</code> , <code>false</code> , <code>no</code> , or <code>off</code> to opt out even on models where the experiment or server configuration would otherwise enable it. The full tool set, hooks, MCP servers, and CLAUDE.md discovery remain enabled
<code>CLAUDE_CODE_SKIP_ANTHROPIC_AWS_AUTH</code>	Skip client-side authentication for Claude Platform on AWS , for gateways that sign requests themselves
<code>CLAUDE_CODE_SKIP_BEDROCK_AUTH</code>	Skip AWS authentication for Bedrock (for example, when using an LLM gateway)
<code>CLAUDE_CODE_SKIP_FOUNDRY_AUTH</code>	Skip Azure authentication for Microsoft Foundry. For a gateway, set the credential in <code>ANTHROPIC_FOUNDRY_API_KEY</code> instead; without an API key this variable leaves the Foundry client unable to send requests
<code>CLAUDE_CODE_SKIP_MANTLE_AUTH</code>	Skip AWS authentication for Bedrock Mantle (for example, when using an LLM gateway)
<code>CLAUDE_CODE_SKIP_PROMPT_HISTORY</code>	Set to <code>1</code> to skip writing prompt history and session transcripts to disk. Sessions started with this variable set do not appear in <code>--resume</code> , <code>--continue</code> , or up-arrow history. Useful for ephemeral scripted sessions
<code>CLAUDE_CODE_SKIP_VERTEX_AUTH</code>	Skip Google authentication for Vertex (for example, when using an LLM gateway)
<code>CLAUDE_CODE_STOP_HOOK_BLOCK_CAP</code>	Maximum number of consecutive times a Stop or SubagentStop hook may block the turn from ending before Claude Code overrides it and ends the turn anyway (default: 8). Set to <code>0</code> to disable the cap. Raise this if your hook legitimately needs more iterations to resolve

Variable	Purpose
CLAUDE_CODE_SUBAGENT_MODEL	See Model configuration
CLAUDE_CODE_SUBPROCESS_ENV_SCRUB	Set to <code>1</code> to strip Anthropic and cloud provider credentials from subprocess environments (Bash tool, hooks, MCP stdio servers). The parent Claude process keeps these credentials for API calls, but child processes cannot read them, reducing exposure to prompt injection attacks that attempt to exfiltrate secrets via shell expansion. On Linux, this also runs Bash subprocesses in an isolated PID namespace so they cannot read host process environments via <code>/proc</code> ; as a side effect, <code>ps</code> , <code>pgrep</code> , and <code>kill</code> cannot see or signal host processes. <code>claude-code-action</code> sets this automatically when <code>allowed_non_write_users</code> is configured
CLAUDE_CODE_SYNC_PLUGIN_INSTALL	Set to <code>1</code> in non-interactive mode (the <code>-p</code> flag) to wait for plugin installation to complete before the first query. Without this, plugins install in the background and may not be available on the first turn. Combine with <code>CLAUDE_CODE_SYNC_PLUGIN_INSTALL_TIMEOUT_MS</code> to bound the wait
CLAUDE_CODE_SYNC_PLUGIN_INSTALL_TIMEOUT_MS	Timeout in milliseconds for synchronous plugin installation. When exceeded, Claude Code proceeds without plugins and logs an error. No default: without this variable, synchronous installation waits until complete
CLAUDE_CODE_SYNC_SKILLS	Set to <code>1</code> to download your enabled claude.ai skills into <code>~/.claude/skills/</code> before the first query and resync every 10 minutes. Applies only in non-interactive mode with the <code>-p</code> flag. Requires claude.ai authentication. Claude Code on the web sessions receive your enabled claude.ai skills automatically; you don't need to set this there
CLAUDE_CODE_SYNC_SKILLS_INSTALL_TIMEOUT_MS	Timeout in milliseconds for a mid-session skills resync when <code>CLAUDE_CODE_SYNC_SKILLS</code> is set (default: 30000). Bounds the download triggered when the host requests a skill reload during the session. When exceeded, the resync stops and remaining downloads continue in the background
CLAUDE_CODE_SYNC_SKILLS_WAIT_TIMEOUT_MS	Timeout in milliseconds for the first query to wait on the initial skills sync when <code>CLAUDE_CODE_SYNC_SKILLS</code> is set (default: 5000). When exceeded, the query proceeds and remaining skill downloads continue in the background
CLAUDE_CODE_SYNTAX_HIGHLIGHT	Set to <code>false</code> to disable syntax highlighting in diff output. Useful when colors interfere with your terminal setup. To also disable highlighting in code blocks and file previews, use the syntaxHighlightingDisabled setting
CLAUDE_CODE_TASK_LIST_ID	Share a task list across sessions. Set the same ID in multiple Claude Code instances to coordinate on a shared task list. See Task list
CLAUDE_CODE_TEAM_NAME	Name of the agent team this teammate belongs to. Set automatically on agent team members
CLAUDE_CODE_TMPDIR	Override the temp directory used for internal temp files. Claude Code appends <code>/claude-{uid}/</code> on Unix or <code>/claude/</code> on Windows to this path. Default: <code>/tmp</code> on macOS, <code>os.tmpdir()</code> on Linux and Windows. As of v2.1.161, on macOS and Linux, sandboxed Bash subprocesses receive a short fallback <code>\$TMPDIR</code> under the system default when your override is a long path, since some tools fail when temp paths get too long. Unsandboxed Bash commands inherit your shell's <code>\$TMPDIR</code> unchanged. Claude Code's own temp files always use your override
CLAUDE_CODE_TMUX_TRUE	Set to <code>1</code> to allow 24-bit truecolor output inside tmux. By default, Claude Code clamps

Variable	Purpose
<code>COLOR</code>	to 256 colors when <code>\$TMUX</code> is set because tmux does not pass through truecolor escape sequences unless configured to. Set this after adding <code>set -ga terminal-overrides ',*:Tc'</code> to your <code>~/.tmux.conf</code> . See Terminal configuration for other tmux settings
<code>CLAUDE_CODE_USE_ANTHROPIC_AWS</code>	Use Claude Platform on AWS
<code>CLAUDE_CODE_USE_BEDROCK</code>	Use Bedrock
<code>CLAUDE_CODE_USE_FOUNDRY</code>	Use Microsoft Foundry
<code>CLAUDE_CODE_USE_MANTLE</code>	Use the Bedrock Mantle endpoint
<code>CLAUDE_CODE_USE_NATIVE_FILE_SEARCH</code>	Set to <code>1</code> to discover custom commands, subagents, and output styles using Node.js file APIs instead of ripgrep. Set this if the bundled ripgrep binary is unavailable or blocked in your environment. Does not affect the Grep or file search tools
<code>CLAUDE_CODE_USE_POWER_SHELL_TOOL</code>	Controls the PowerShell tool. On Windows without Git Bash, the tool is enabled automatically; set to <code>0</code> to disable it. On Windows with Git Bash installed, the tool is rolling out progressively: set to <code>1</code> to opt in or <code>0</code> to opt out. On Linux, macOS, and WSL, set to <code>1</code> to enable it, which requires <code>pwsh</code> on your <code>PATH</code> . When enabled on Windows, Claude can run PowerShell commands natively instead of routing through Git Bash. See PowerShell tool
<code>CLAUDE_CODE_USE_VERTEX</code>	Use Vertex
<code>CLAUDE_CONFIG_DIR</code>	Override the configuration directory (default: <code>~/.claude</code>). All settings, credentials, session history, and plugins are stored under this path. Useful for running multiple accounts side by side: for example, <code>alias claude-work='CLAUDE_CONFIG_DIR=~/.claude-work claude'</code>
<code>CLAUDE_EFFORT</code>	Set automatically in Bash tool subprocesses and hook commands to the active effort level for the turn: <code>low</code> , <code>medium</code> , <code>high</code> , <code>xhigh</code> , or <code>max</code> . Ultracode is not a distinct level and reports as <code>xhigh</code> . Matches the <code>effort.level</code> field passed to hooks . Only set when the current model supports the effort parameter
<code>CLAUDE_ENABLE_BYTE_WATCHDOG</code>	Set to <code>1</code> to force-enable the byte-level streaming idle watchdog, or set to <code>0</code> to force-disable it. When unset, the watchdog is enabled by default for direct Anthropic API and Claude Platform on AWS connections. The byte watchdog aborts a connection when no bytes arrive on the wire for 180 seconds by default on direct Anthropic API connections, 300 seconds on Claude Platform on AWS and when enabled on Bedrock, or for the value of <code>CLAUDE_STREAM_IDLE_TIMEOUT_MS</code> when that is set, which is clamped to a minimum of 5 minutes, independent of the event-level watchdog
<code>CLAUDE_ENABLE_BYTE_WATCHDOG_BEDROCK</code>	Set to <code>1</code> to enable the byte-level streaming idle watchdog on Amazon Bedrock <code>vnd.amazon.eventstream</code> responses. Off by default. Configure the timeout with <code>CLAUDE_STREAM_IDLE_TIMEOUT_MS</code>
<code>CLAUDE_ENABLE_STREAM_WATCHDOG</code>	Set to <code>1</code> to force-enable the event-level streaming idle watchdog, or set to <code>0</code> to force-disable it. When unset, the default is server-controlled on the direct Anthropic API

Variable	Purpose
	and off on other providers. As of v2.1.169, providers other than the direct Anthropic API and Claude Platform on AWS also have a default-on 5-minute body idle timeout independent of this variable; see <code>API_FORCE_IDLE_TIMEOUT</code> . On Bedrock, you can also enable the independent byte-level watchdog with <code>CLAUDE_ENABLE_BYTE_WATCHDOG_BEDROCK</code> ; the two run together when both are set. Configure the timeout with <code>CLAUDE_STREAM_IDLE_TIMEOUT_MS</code>
<code>CLAUDE_ENV_FILE</code>	Path to a shell script whose contents Claude Code runs before each Bash command in the same shell process, so exports in the file are visible to the command. Use to persist virtualenv or conda activation across commands. Also populated dynamically by <code>SessionStart</code> , <code>Setup</code> , <code>CwdChanged</code> , and <code>FileChanged</code> hooks
<code>CLAUDE_REMOTE_CONTROL_SESSION_NAME_PREFIX</code> <code>X</code>	Prefix for auto-generated <code>Remote Control</code> session names when no explicit name is provided. Defaults to your machine's hostname, producing names like <code>myhost-graceful-unicorn</code> . The <code>--remote-control-session-name-prefix</code> CLI flag sets the same value for a single invocation
<code>CLAUDE_STREAM_IDLE_TIMEOUT_MS</code>	Timeout in milliseconds before the streaming idle watchdog closes a stalled connection. When you set this variable explicitly, the minimum is <code>300000</code> (5 minutes); lower values are silently clamped to absorb extended thinking pauses and proxy buffering. When unset, the event-level watchdog defaults to 300 seconds and the byte-level watchdog defaults to 180 seconds on direct Anthropic API connections (300 seconds on Claude Platform on AWS and other providers). The unset 180-second byte-watchdog default is a separate value and is not subject to the 5-minute clamp. For the event-level watchdog on third-party providers, requires <code>CLAUDE_ENABLE_STREAM_WATCHDOG=1</code> ; the body idle timeout described under <code>API_FORCE_IDLE_TIMEOUT</code> applies independently. On Bedrock, also applies when <code>CLAUDE_ENABLE_BYTE_WATCHDOG_BEDROCK=1</code>
<code>DEBUG</code>	Set to <code>1</code> to enable debug mode, equivalent to launching with <code>--debug</code> . Debug logs are written to <code>~/.claude/debug/<session-id>.txt</code> , or to the path set by <code>CLAUDE_CODE_DEBUG_LOGS_DIR</code> . Only the truthy values <code>1</code> , <code>true</code> , <code>yes</code> , and <code>on</code> enable debug mode, so namespace patterns like <code>DEBUG=express:*</code> set for other tools do not trigger it
<code>DISABLE_AUTOUPDATER</code>	Set to <code>1</code> to disable automatic background updates. Manual <code>claude update</code> still works. Use <code>DISABLE_UPDATES</code> to block both
<code>DISABLE_AUTO_COMPACT</code> <code>T</code>	Set to <code>1</code> to disable automatic compaction when approaching the context limit. The manual <code>/compact</code> command remains available. Use when you want explicit control over when compaction occurs
<code>DISABLE_COMPACT</code>	Set to <code>1</code> to disable all compaction: both automatic compaction and the manual <code>/compact</code> command
<code>DISABLE_COST_WARNINGS</code>	Set to <code>1</code> to disable cost warning messages
<code>DISABLE_DOCTOR_COMMAND</code>	Set to <code>1</code> to hide the <code>/doctor</code> command. Useful for managed deployments where users should not run installation diagnostics
<code>DISABLE_ERROR_REPORTING</code>	Set to <code>1</code> to opt out of Sentry error reporting
<code>DISABLE_EXTRA_USAGE_C</code>	Set to <code>1</code> to hide the <code>/usage-credits</code> command that lets users purchase additional

Variable	Purpose
OMMAND	usage beyond rate limits
DISABLE_FEEDBACK_COMMAND AND	Set to <code>1</code> to disable the <code>/feedback</code> command. The older name <code>DISABLE_BUG_COMMAND</code> is also accepted
DISABLE_GROWTHBOOK	Set to <code>1</code> to disable GrowthBook feature-flag fetching and use code defaults for every flag. Telemetry event logging stays on unless <code>DISABLE_TELEMETRY</code> is also set
DISABLE_INSTALLATION_CHECKS	Set to <code>1</code> to disable installation warnings. Use only when manually managing the installation location, as this can mask issues with standard installations
DISABLE_INSTALL_GITHUB_APP_COMMAND	Set to <code>1</code> to hide the <code>/install-github-app</code> command. Already hidden when using third-party providers (Bedrock, Vertex, or Foundry)
DISABLE_INTERLEAVED_THINKING	Set to <code>1</code> to prevent sending the interleaved-thinking beta header. Useful when your LLM gateway or provider does not support <u>interleaved thinking</u>
DISABLE_LOGIN_COMMAND	Set to <code>1</code> to hide the <code>/login</code> command. Useful when authentication is handled externally via API keys or <code>apiKeyHelper</code>
DISABLE_LOGOUT_COMMAND	Set to <code>1</code> to hide the <code>/logout</code> command
DISABLE_PROMPT_CACHING	Set to <code>1</code> to disable <u>prompt caching</u> for all models (takes precedence over per-model settings)
DISABLE_PROMPT_CACHING_FABLE	Set to <code>1</code> to disable prompt caching for Fable models
DISABLE_PROMPT_CACHING_HAIKU	Set to <code>1</code> to disable prompt caching for Haiku models
DISABLE_PROMPT_CACHING_OPUS	Set to <code>1</code> to disable prompt caching for Opus models
DISABLE_PROMPT_CACHING_SONNET	Set to <code>1</code> to disable prompt caching for Sonnet models
DISABLE_TELEMETRY	Set to <code>1</code> to opt out of telemetry. Telemetry events do not include user data like code, file paths, or bash commands. Also disables feature-flag fetching with the same effect as <code>DISABLE_GROWTHBOOK</code> , so some flagged features may be unavailable
DISABLE_UPDATES	Set to <code>1</code> to block all updates including manual <code>claude update</code> and <code>claude install</code> . Stricter than <code>DISABLE_AUTOUPDATER</code> . Use when distributing Claude Code through your own channels and users should not self-update
DISABLE_UPGRADE_COMMAND	Set to <code>1</code> to hide the <code>/upgrade</code> command
DO_NOT_TRACK	Set to <code>1</code> to opt out of telemetry. Equivalent to setting <code>DISABLE_TELEMETRY</code> . Claude Code honors this as the cross-tool convention recognized by many developer CLIs
ENABLE_CLAUDEAI_MCP_SERVERS	Set to <code>false</code> to disable <u>claude.ai MCP servers</u> in Claude Code. Enabled by default for logged-in users. To disable per-project or per-org, set <u>disableClaudeAiConnectors</u> in settings instead

Variable	Purpose
<code>ENABLE_PROMPT_CACHING_1H</code>	Set to <code>1</code> to request a 1-hour prompt cache TTL instead of the default 5 minutes. Intended for API key, Bedrock , Vertex , Foundry , and Claude Platform on AWS users. Subscription users within included usage receive 1-hour TTL automatically. 1-hour cache writes are billed at a higher rate
<code>ENABLE_PROMPT_CACHING_1H_BEDROCK</code>	Deprecated. Use <code>ENABLE_PROMPT_CACHING_1H</code> instead
<code>ENABLE_TOOL_SEARCH</code>	Controls MCP tool search . Unset: all MCP tools deferred by default, but loaded upfront on Vertex AI or when <code>ANTHROPIC_BASE_URL</code> points to a non-first-party host. Values: <code>true</code> (always defer and send the beta header, requests fail on Vertex AI models earlier than Sonnet 4.5 or Opus 4.5, or on proxies that do not support <code>tool_reference</code>), <code>auto</code> (threshold mode: load upfront if tools fit within 10% of context), <code>auto:N</code> (custom threshold, e.g., <code>auto:5</code> for 5%), <code>false</code> (load all upfront)
<code>FALLBACK_FOR_ALL_PRIMARY_MODELS</code>	Set to any non-empty value to make all models, not only Opus, stop retrying with a repeated-overload error when no fallback model is configured. As of v2.1.160, a configured fallback model chain triggers on repeated overload errors for any primary model, so this variable does not affect switching to a fallback model
<code>FORCE_AUTOUPDATE_PLUGINS</code>	Set to <code>1</code> to force plugin auto-updates even when the main auto-updater is disabled via <code>DISABLE_AUTOUPDATER</code>
<code>FORCE_PROMPT_CACHING_5M</code>	Set to <code>1</code> to force the 5-minute prompt cache TTL even when 1-hour TTL would otherwise apply. Overrides <code>ENABLE_PROMPT_CACHING_1H</code>
<code>HTTP_PROXY</code>	Specify HTTP proxy server for network connections
<code>HTTPS_PROXY</code>	Specify HTTPS proxy server for network connections
<code>IS_DEMO</code>	Set to <code>1</code> to enable demo mode: hides your email and organization name from the header and <code>/status</code> output, and skips onboarding. Useful when streaming or recording a session
<code>MAX_MCP_OUTPUT_TOKENS</code>	Maximum number of tokens allowed in MCP tool responses. Claude Code displays a warning when output exceeds 10,000 tokens. Tools that declare <code>anthropic/maxResultSizeChars</code> use that character limit for text content instead, but image content from those tools is still subject to this variable (default: 25000)
<code>MAX_STRUCTURED_OUTPUT_RETRIES</code>	Number of times to retry when the model's response fails validation against the <code>--json-schema</code> in non-interactive mode (the <code>-p</code> flag). Defaults to 5
<code>MAX_THINKING_TOKENS</code>	Override the extended thinking token budget. The ceiling is the model's max output tokens minus one. Set to <code>0</code> to disable thinking on the Anthropic API except on Fable 5, which cannot have thinking turned off. On third-party providers , <code>0</code> omits the <code>thinking</code> parameter instead, and models with adaptive reasoning may still think. For nonzero values on adaptive reasoning models, the budget is ignored unless adaptive reasoning is disabled via <code>CLAUDE_CODE_DISABLE_ADAPTIVE_THINKING</code>
<code>MCP_CLIENT_SECRET</code>	OAuth client secret for MCP servers that require pre-configured credentials . Avoids the interactive prompt when adding a server with <code>--client-secret</code>
<code>MCP_CONNECTION_NONBLOCKING</code>	Controls whether startup waits for MCP servers to connect before the first query. As of Claude Code v2.1.142, MCP startup is non-blocking by default: servers connect in the

Variable	Purpose
	background and their tools become available as they finish. Set to <code>0</code> to restore the blocking 5-second connection wait. Servers configured with <code>alwaysLoad: true</code> still block startup regardless, since their tools must be present when the first prompt is built
<code>MCP_CONNECT_TIMEOUT_MS</code>	How long blocking MCP startup waits, in milliseconds, for the connection batch before snapshotting the tool list (default: 5000). Applies when <code>MCP_CONNECTION_NONBLOCKING=0</code> or for servers marked <code>alwaysLoad: true</code> . Servers still pending at the deadline keep connecting in the background but won't appear until the next query. Distinct from <code>MCP_TIMEOUT</code> , which bounds an individual server's connect attempt
<code>MCP_OAUTH_CALLBACK_PORT</code>	Fixed port for the OAuth redirect callback, as an alternative to <code>--callback-port</code> when adding an MCP server with <u>pre-configured credentials</u>
<code>MCP_REMOTE_SERVER_CONNECTION_BATCH_SIZE</code>	Maximum number of remote MCP servers (HTTP/SSE) to connect in parallel during startup (default: 20)
<code>MCP_SERVER_CONNECTION_BATCH_SIZE</code>	Maximum number of local MCP servers (stdio) to connect in parallel during startup (default: 3)
<code>MCP_TIMEOUT</code>	Timeout in milliseconds for MCP server startup (default: 30000, or 30 seconds)
<code>MCP_TOOL_TIMEOUT</code>	Timeout in milliseconds for MCP tool execution (default: 100000000, about 28 hours). A per-server <code>timeout</code> field in <code>.mcp.json</code> overrides this for that server. For the env variable, values below 1000 are floored to one second; for the per-server field, values below 1000 are ignored
<code>NO_PROXY</code>	List of domains and IPs to which requests will be directly issued, bypassing proxy
<code>OTEL_LOG_RAW_API_BODIES</code>	Emit Anthropic Messages API request and response JSON as <code>api_request_body</code> / <code>api_response_body</code> log events. Set to <code>1</code> for inline bodies truncated at 60 KB, or <code>file:<dir></code> to write untruncated bodies to disk and emit a <code>body_ref</code> path instead. Disabled by default; bodies include the entire conversation history. See <u>Monitoring</u>
<code>OTEL_LOG_TOOL_CONTENT</code>	Set to <code>1</code> to include tool input and output content in OpenTelemetry span events. Disabled by default to protect sensitive data. See <u>Monitoring</u>
<code>OTEL_LOG_TOOL_DETAILS</code>	Set to <code>1</code> to include tool input arguments, MCP server names, raw error strings on tool failures, the refusal <code>category</code> on <code>api_refusal</code> events, and other tool details in OpenTelemetry traces and logs. Disabled by default to protect PII. See <u>Monitoring</u>
<code>OTEL_LOG_USER_PROMPTS</code>	Set to <code>1</code> to include user prompt text in OpenTelemetry traces and logs. Disabled by default (prompts are redacted). See <u>Monitoring</u>
<code>OTEL_METRICS_INCLUDE_ACCOUNT_UUID</code>	Set to <code>false</code> to exclude account UUID from metrics attributes (default: included). See <u>Monitoring</u>
<code>OTEL_METRICS_INCLUDE_ENTRYPOINT</code>	Set to <code>true</code> to include the session entrypoint in metrics attributes (default: excluded). Added in v2.1.152. See <u>Monitoring</u>
<code>OTEL_METRICS_INCLUDE_RESOURCE_ATTRIBUTES</code>	As of v2.1.161, Claude Code attaches <code>OTEL_RESOURCE_ATTRIBUTES</code> keys to metric datapoint labels. Set to <code>false</code> to exclude them (default: included). See <u>Monitoring</u>

Variable	Purpose
OTEL_METRICS_INCLUDE_SESSION_ID	Set to <code>false</code> to exclude session ID from metrics attributes (default: included). See Monitoring
OTEL_METRICS_INCLUDE_VERSION	Set to <code>true</code> to include Claude Code version in metrics attributes (default: excluded). See Monitoring
SLASH_COMMAND_TOOL_CHARACTER_BUDGET	Override the character budget for skill metadata shown to the Skill tool . The budget scales dynamically at 1% of the context window, with a fallback of 8,000 characters. Legacy name kept for backwards compatibility
TASK_MAX_OUTPUT_LENGTH	Maximum number of characters in subagent output before truncation (default: 32000, maximum: 160000). When truncated, the full output is saved to disk and the path is included in the truncated response
USE_BUILTIN_RIPGREP	Set to <code>0</code> to use system-installed <code>rg</code> instead of <code>rg</code> included with Claude Code
VERTEX_REGION_CLAUDE_3_5_HAIKU	Override region for Claude 3.5 Haiku when using Vertex AI
VERTEX_REGION_CLAUDE_3_5_SONNET	Override region for Claude 3.5 Sonnet when using Vertex AI
VERTEX_REGION_CLAUDE_3_7_SONNET	Override region for Claude 3.7 Sonnet when using Vertex AI
VERTEX_REGION_CLAUDE_4_0_OPUS	Override region for Claude 4.0 Opus when using Vertex AI
VERTEX_REGION_CLAUDE_4_0_SONNET	Override region for Claude 4.0 Sonnet when using Vertex AI
VERTEX_REGION_CLAUDE_4_1_OPUS	Override region for Claude 4.1 Opus when using Vertex AI
VERTEX_REGION_CLAUDE_4_5_OPUS	Override region for Claude Opus 4.5 when using Vertex AI
VERTEX_REGION_CLAUDE_4_5_SONNET	Override region for Claude Sonnet 4.5 when using Vertex AI
VERTEX_REGION_CLAUDE_4_6_OPUS	Override region for Claude Opus 4.6 when using Vertex AI
VERTEX_REGION_CLAUDE_4_6_SONNET	Override region for Claude Sonnet 4.6 when using Vertex AI
VERTEX_REGION_CLAUDE_4_7_OPUS	Override region for Claude Opus 4.7 when using Vertex AI. Added in v2.1.111
VERTEX_REGION_CLAUDE_4_8_OPUS	Override region for Claude Opus 4.8 when using Vertex AI. Added in v2.1.154
VERTEX_REGION_CLAUDE_FABLE_5	Override region for Claude Fable 5 when using Vertex AI. Added in v2.1.170

Variable	Purpose
VERTEX_REGION_CLAUDE_	Override region for Claude Haiku 4.5 when using Vertex AI
HAIKU_4_5	

Standard OpenTelemetry exporter variables (`OTEL_METRICS_EXPORTER` , `OTEL_LOGS_EXPORTER` , `OTEL_EXPORTER_OTLP_ENDPOINT` , `OTEL_EXPORTER_OTLP_PROTOCOL` , `OTEL_EXPORTER_OTLP_HEADERS` , `OTEL_METRIC_EXPORT_INTERVAL` , `OTEL_RESOURCE_ATTRIBUTES` , and signal-specific variants) are also supported. See **Monitoring** for configuration details.

See also

- **Settings**: all `settings.json` configuration, including the `env` key
- **CLI reference**: launch-time flags
- **Network configuration**: proxy and TLS setup
- **Monitoring**: OpenTelemetry configuration