

WHAT'S NEW

Week 16 · April 13–17, 2026

Claude Opus 4.7 with the new `xhigh` effort level, Routines on Claude Code on the web, mobile push notifications that ping your phone when Claude needs you, a `/usage` breakdown that shows what's driving your limits, and native binaries replacing the bundled JavaScript.

Releases **v2.1.105 → v2.1.113**

5 features · April 13–17

Claude Opus 4.7 new model

Anthropic's strongest coding model yet is now the default on Max and Team Premium, and available everywhere else from `/model`. It adds a new `xhigh` effort level that sits between `high` and `max`: best results for most coding and agentic tasks, applied as the default the first time you switch to 4.7. `/effort` now opens an interactive arrow-key slider when you call it without arguments, so you can dial intelligence against speed without remembering the level names.

Switch model and effort in one go:

Claude Code


```
> /model opus
> /effort xhigh
```

Model config: effort levels →

Routines web

Templated cloud agents that fire on a schedule, a GitHub event, or an API call. Define a

routine once on Claude Code on the web with a prompt, the repos it can touch, and the connectors it needs, then let PR-opened, release-published, or your own webhook trigger it without your machine running. The trigger picker now covers GitHub events with optional filters and gives every routine a tokened `/fire` endpoint for external systems.



Create one from the web UI, or scaffold from your terminal:

Claude Code

```
> /schedule daily PR review at 9am
```

[Routines guide](#) →

/usage breakdown CLI

More visibility into where your Claude Code usage goes. `/usage` now shows what's driving your limits: parallel sessions, subagents, cache misses, and long context, each with a percentage of your last 24 hours and a tip to optimize it. Press `d` or `w` to switch between day and week views.

Try `/usage`

```

~/project - claude code

What's contributing to your limits usage?
Approximate, based on local sessions on this machine – does not include other devices or
claude.ai

Last 24h · these are independent characteristics of your usage, not a breakdown

98% of your usage was while 4+ sessions ran in parallel
  All sessions share one limit. If you don't need them all at once, queueing uses it more
  evenly.

90% of your usage came from sessions that ran 3+ subagents
  Each subagent runs its own requests, using up your limits faster. Reserve parallel subagent
  use for when you really need it.

75% of your usage was at >150k context
  Longer sessions are more expensive even when cached. /compact mid-task, /clear when switching
  to new tasks.

19% of your usage hit a >100k-token cache miss
  Uncached input is expensive, and often happens when sending a message to a session that has
  gone idle. /compact before stepping away keeps the cold-start small.

12% of your usage came from sessions active for 8+ hours
  These are often background/loop sessions. Continuous usage can add up quickly so make sure it
  is intentional.

d to day · w to week

```

Run it any time:

Claude Code


> `/usage`

[Commands reference](#) →

Mobile push notifications mobile

With [Remote Control](#) connected, Claude can send a push notification to your phone

when a long task finishes or it needs a decision to keep going. Turn it on with “Push when Claude decides” in `/config`, or ask for one in your prompt. Useful when you kick off a long agent run and want to step away from the terminal.



Ask Claude to ping you when it's done:

Claude Code

```
> notify me when the tests pass
```

[Remote Control: mobile push notifications →](#)

Native binaries v2.1.113

The `claude` CLI now spawns a native per-platform binary instead of bundled JavaScript, so the installed `claude` command no longer invokes Node. The `npm` package pulls the right binary in through an optional dependency such as `@anthropic-ai/claude-code-darwin-arm64`, so your install command doesn't change. The standalone installer already shipped this binary; `npm` now matches it.

Upgrade and check what you're running:

```
claude update
claude --version
```

[Setup guide →](#)

Other wins

- + New `/ultrareview`: comprehensive code review in the cloud using parallel multi-agent analysis and an adversarial critique pass. Run it bare to review your current branch, or `/ultrareview <PR#>` for a specific PR
- + **Session recap** shows a one-line summary of what happened while you were away; run `/recap` on demand or turn it off from `/config`
- + Plugins can ship background watchers via a top-level `monitors` manifest key that auto-arms at session start or on skill invoke
- + `/fewer-permission-prompts` scans your transcripts for common read-only Bash and MCP calls and proposes an allowlist for `.claude/settings.json`
- + `PreCompact` hooks can block compaction by exiting with code 2 or returning `{"decision": "block"}`
- + `sandbox.network.deniedDomains` setting carves specific domains out of a broader `allowedDomains` wildcard
- + Hardened Bash permissions: deny rules now match through `env` / `sudo` / `watch` wrappers, and `Bash(find:*)` allow rules no longer auto-approve `-exec` or `-delete`
- + **Auto mode** is now available for Max subscribers on Opus 4.7, and the `-enable-auto-mode` flag is no longer required
- + New `/tui` command and `tui` setting switch between classic and flicker-free rendering mid-conversation; focus view moved from `Ctrl+O` to its own `/focus` command
- + "Auto (match terminal)" option in `/theme` follows your terminal's dark/light mode
- + Claude can now discover and run built-in commands like `/init`, `/review`, and `/security-review` via the Skill tool
- + `ENABLE_PROMPT_CACHING_1H` opts API key, Bedrock, Vertex, and Foundry users into 1-hour prompt cache TTL
- + `/undo` is now an alias for `/rewind`, and `/proactive` is an alias for `/loop`

[Full changelog for v2.1.105–v2.1.113 →](#)